

Process Management / Shell

Programmer Interface

- ❑ **Kernel:** Everything below the system-call interface and above the physical hardware
 - Provides file system, CPU scheduling, memory management, and other OS functions through system calls
- ❑ **Systems programs:** Use the kernel-supported system calls to provide useful functions, such as compilation and file manipulation

UNIX Layer Structure

(the users)		
shells and commands compilers and interpreters system libraries		
<i>system-call interface to the kernel</i>		
signals terminal handling character I/O system terminal drivers	file system swapping block I/O system disk and tape drivers	CPU scheduling page replacement demand paging virtual memory
<i>kernel interface to the hardware</i>		
terminal controllers terminals	device controllers disks and tapes	memory controllers physical memory

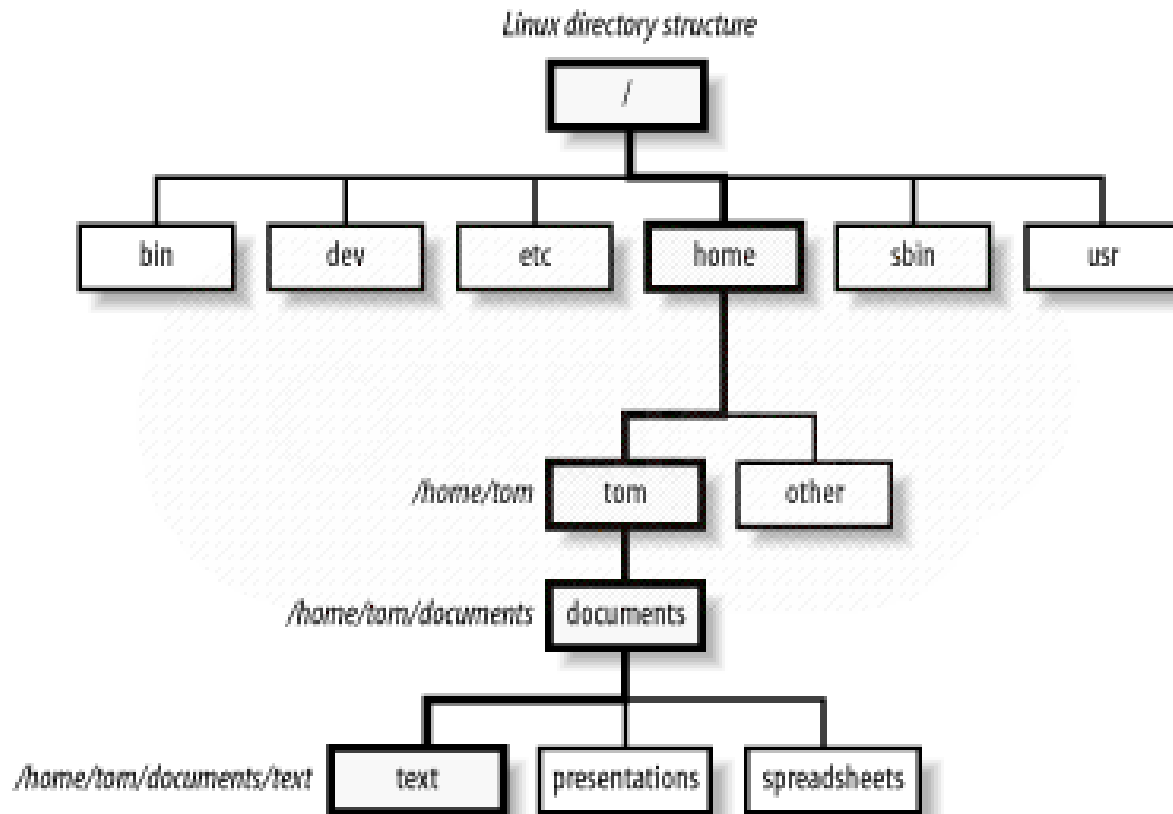
System Calls

- ❑ System calls define the programmer interface to UNIX
- ❑ The set of systems programs commonly available defines the user interface
- ❑ The programmer and user interface define the context that the kernel must support
- ❑ Roughly three categories of system calls in UNIX
 - File manipulation (same system calls also support device manipulation)
 - Process control
 - Information manipulation

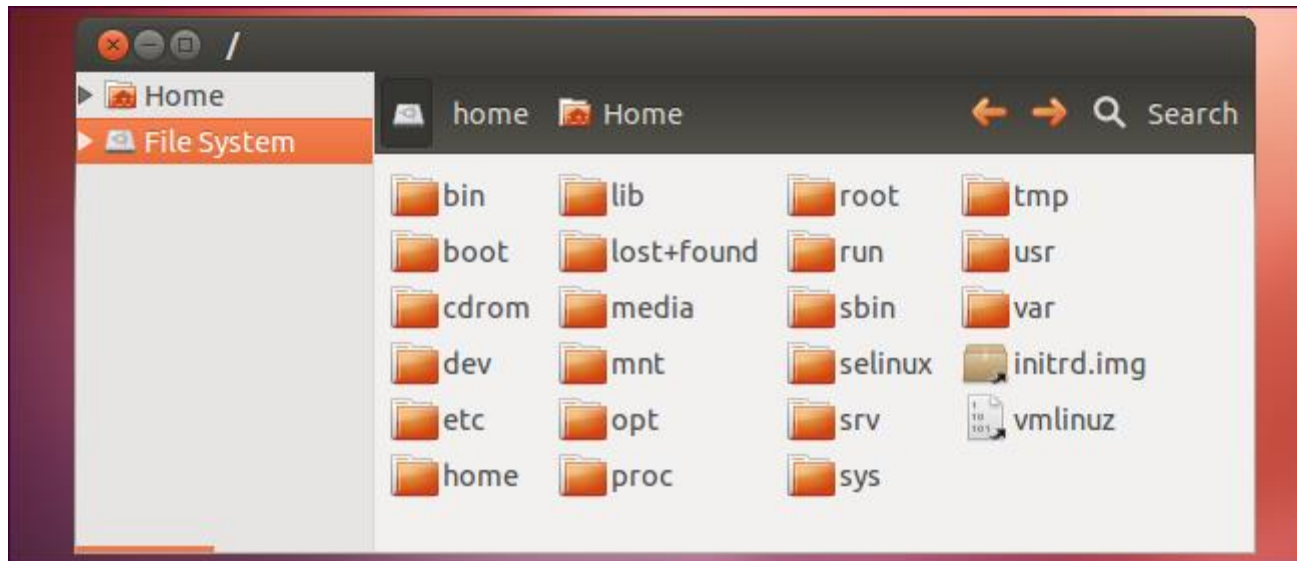
File Manipulation

- ❑ A file is a sequence of bytes; the kernel does not impose a structure on files
- ❑ Files are organized in tree-structured **directories**
- ❑ Directories are files that contain information on how to find other files
- ❑ **Path name:** identifies a file by specifying a path through the directory structure to the file
 - Absolute path names start at root of file system
 - Relative path names start at the current directory
- ❑ **System calls for basic file manipulation:** `creat`, `open`, `read`, `write`, `close`, `unlink`, `trunc`

Typical UNIX Directory Structure



Ubuntu Directory Structure



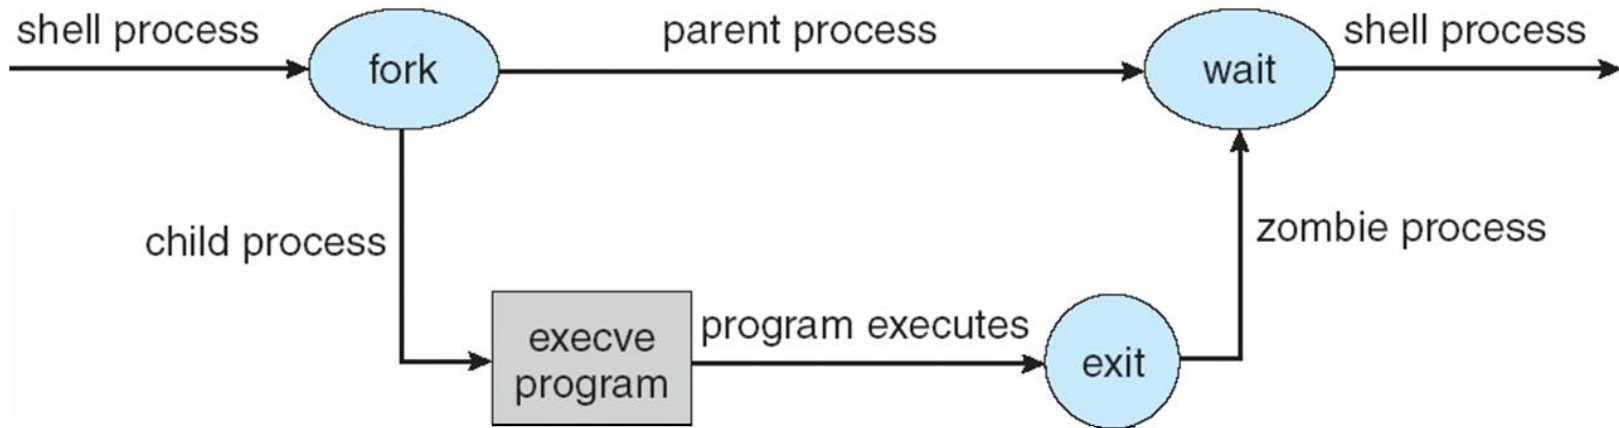
Process Control

- ❑ A process is a program in execution
- ❑ Processes are identified by their process identifier, an integer
- ❑ Process control system calls
 - `fork` creates a new process
 - `execve` is used after a `fork` to replace one of the two processes's virtual memory space with a new program
 - `exit` terminates a process
 - A parent may `wait` for a child process to terminate; `wait` provides the process id of a terminated child so that the parent can tell which child terminated

Process Control (cont.)

- ❑ A **zombie** process results when the parent of a **defunct** child process exits before the terminated child.
 - When a process dies it is not all removed from memory immediately.
 - The process status becomes `EXIT_ZOMBIE` and the process's parent is notified that its child process has died.
 - Parent process is supposed to execute the `wait()` system call to read the dead process's exit status.
 - Zombie children processes will stick around in memory until they are cleaned up.
 - The **top** command will display information about **running**, **sleeping**, **stopped**, and **zombie** processes in Linux.

Illustration of Process Control Calls



Process Control (Cont.)

- ❑ Processes communicate via pipes; queues of bytes between two processes that are accessed by a file descriptor.
- ❑ All user processes are descendants of one original process, *init*
- ❑ *init* forks a *getty* (get tty) process: initializes terminal line parameters and passes the user's *login name* to *login*.
 - *login* sets the numeric *user identifier* of the process to that of the user.
 - executes a **shell** which forks subprocesses for user commands.

Information Manipulation

- ❑ System calls to set and return an interval timer:

`getitimer/setitimer`

- ❑ Calls to set and return the current time:

`gettimeofday/settimeofday`

- ❑ Processes can ask for

- their process identifier: `getpid`

- their group identifier: `getgid`

- the name of the machine on which they are executing:

`gethostname`

Library Routines

- ❑ The system-call interface to UNIX is supported and augmented by a large collection of library routines.
- ❑ Header files provide the definition of complex data structures used in system calls.
- ❑ Additional library support is provided for mathematical functions, network access, data conversion, etc.

User Interface

- ❑ Programmers and users mainly deal with already existing systems programs: the needed system calls are embedded within the program and do not need to be obvious to the user.
- ❑ The most common systems programs are file or directory oriented.
 - Directory: `mkdir`, `rmdir`, `cd`, `pwd`
 - File: `ls`, `cp`, `mv`, `rm`
- ❑ Other programs relate to editors (e.g., `emacs`, `vi`) text formatters (e.g., `troff`, `TeX`), and other activities.

Standard I/O

- ❑ Most processes expect three file descriptors to be open when they start:
 - *standard input* - program can read what the user types.
 - *standard output* - program can send output to user's screen.
 - *standard error* - error output.
- ❑ Most programs can also accept a file (rather than a terminal) for standard input and standard output.
- ❑ The common shells have a simple syntax for changing what files are open for the standard I/O streams of a process — *I/O redirection*.

Standard I/O Redirection

Command	Meaning of command
% ls > filea	Direct output of ls to file filea.
% pr < filea > fileb	Input from filea and output to fileb.
% lpr < fileb	Input from fileb.
% ./fork1 > errs &	Save output in a file.

The Shell

- ❑ The **command interpreter** accepts commands from the user. It surrounds the kernel.
- ❑ The **shell** refers to the command interpreter being executed.
- ❑ Commands that we type in a Unix-based system such as "ls" or "ps" are not actually considered part of the operating system.
- ❑ The commands make use of **system calls**. They are executable binary object files.
- ❑ A system call allows a command to request a service from the operating system.
- ❑ A shell is a process.
 - A **process** e.g., a program being executed.

The Shell

- ❑ There are different shells that you can use in a Unix-based system including:
 - bourne shell
 - C shell
 - bash shell
 - tcsh shell
 - and many more
- ❑ A list of various shells may be found at:
 - https://en.wikipedia.org/wiki/Comparison_of_command_shells
 - https://en.wikipedia.org/wiki/Unix_shell

The Shell

- ❑ When a user logs in, a shell is started up.
- ❑ The shell has the terminal as standard input and display as standard output
- ❑ The shell starts out by typing the prompt, a character such as a dollar sign or percentage sign e.g.,

jackson\$

- ❑ User enters a command e.g.,

jackson\$ date

The Shell

- ❑ The user can specify that standard output be redirected to a file e.g.,

`date >file`

- ❑ Standard input can also be redirected e.g.,

`sort <file1 >file2`

- ❑ The output of one program can be used as the input to another program:

`cat file1 file2 file3 | sort >/dev/lp &`

High-Level View of Shell Code

```
while (1)
{
    Get a line from the user.
    Execute command found in line.
}
```

Details Not Highlighted


- Making sure that the line from the user is correct.
- How is shell termination handled?

- The execution of a command is done by a separate process (**child process**) from the shell process.
- For simple commands, the shell process waits for the child process to terminate so that it can print the prompt.
- If a child process is put in the background (using & the ampersand symbol) then the shell process can continue without waiting for the child process to terminate.

Unix Process Creation

- ❑ The Unix system call for process creation is called **fork()**.
- ❑ The **fork** system call creates a child process that is a duplicate of the parent.
 - Child inherits state from parent process.
 - Same program instructions, variables have the same values, same position in the code.
 - Parent and child have separate copies of that state.
 - Child has the same open file descriptors from the parent.
 - Parent and child file descriptors point to a common entry in the system open file table.

Memory Image of a Unix Process



- ❑ Processes have three types of memory segments:
 - Text: program code.
 - Data:
 - Statically declared variables.
 - Heap
 - Areas allocated by `malloc()` or `new` (heap).
 - Stack
 - Automatic variables.
 - Function and system calls.
- ❑ Invoking the same program multiple times results in the creation of multiple distinct address spaces.

Unix Process Management


- ❑ Unix fork()
 - ❑ Creates a new address space.
 - ❑ Copies text, data, and stack into new address space.
 - ❑ Provides child with access to opened files.
- ❑ Unix exec()
 - ❑ Allows a child to run a new program.
 - ❑ There is no system call or library function with the name exec. There is a family of functions, and we refer to them as exec.
- ❑ Unix wait()
 - ❑ Allow a parent to wait for a child to terminate.

Why Create a New Process?

- ❑ Run a new program:
 - E.g., shell executing a program entered at command line.
 - Or, even running an entire pipeline of commands
 - Such as `"wc -l * | sort | uniq -c | sort -nr"`
- ❑ Run a new thread of control for the same program:
 - E.g., a Web server handling a new Web request
 - While continuing to allow more requests to arrive.
 - Essentially time sharing the computer.
- ❑ Underlying mechanism:
 - A process executes `fork()` to create a child process.
 - (Optionally) child does `exec()` of a new program.

Creating a New Process

- ❑ Cloning an existing process.
 - Parent process creates a new child process.
 - The two processes then run concurrently.
- ❑ Child process inherits state from parent.
 - Identical (but separate) copy of virtual address space.
 - Copy of the parent's open file descriptors.
 - Parent and child share access to open files.
- ❑ Child then runs independently.
 - Executing independently, including invoking a new program.
 - Reading and writing its own address space.



Fork System-Level Function

- ❑ `fork()` is called once.
 - But returns twice, once in each process.
- ❑ Telling which process is which.
 - Parent: `fork()` returns the child's process ID.
 - Child: `fork()` returns 0.

```
pid = fork();  
if (pid > 0) {  
    /* in parent */ ...  
} else if (pid == 0) {  
    /* in child */ ...  
} else {  
    /* Error */ ...  
}
```

Fork and Process State

❑ Inherited

- User and group IDs
- Signal handling settings
- stdio
- File pointers
- Root directory
- File mode creation mask
- Resource limits
- Controlling terminal
- All machine register states
- Control register(s)
- ...

❑ Separate in child

- Process ID
- Address space (memory)
- File descriptors
- Parent process ID
- Pending signals
- Time signal reset times
- ...

Example: What Output?

```
int main(void)
{ pid_t pid;
  int x = 1;

  pid = fork();
  if( pid > 0 ) {
    printf("parent: x = %d\n", --x); exit(0);
  } else if( pid < 0 ) {
    printf("Error\n"); exit(0);
  } else {
    printf("child: x = %d\n", ++x);
    exit(0);
  }
}
```

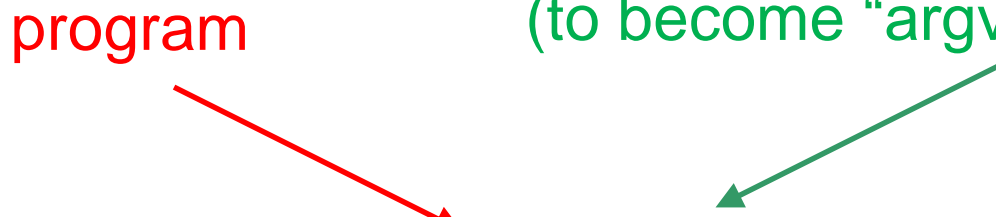
Executing a New Program

- ❑ `fork()` copies the state of the parent process.
 - Child continues running the parent program
 - ... with a copy of the process memory and registers.
- ❑ Need a way to invoke a new program.
 - In the context of the newly-created child process
- ❑ Example:

NULL-terminated array

Contains command-line arguments
(to become “argv[]” of `ls`)

program



```
execvp("ls", argv);  
fprintf(stderr, "exec failed\n");  
exit(EXIT_FAILURE);
```

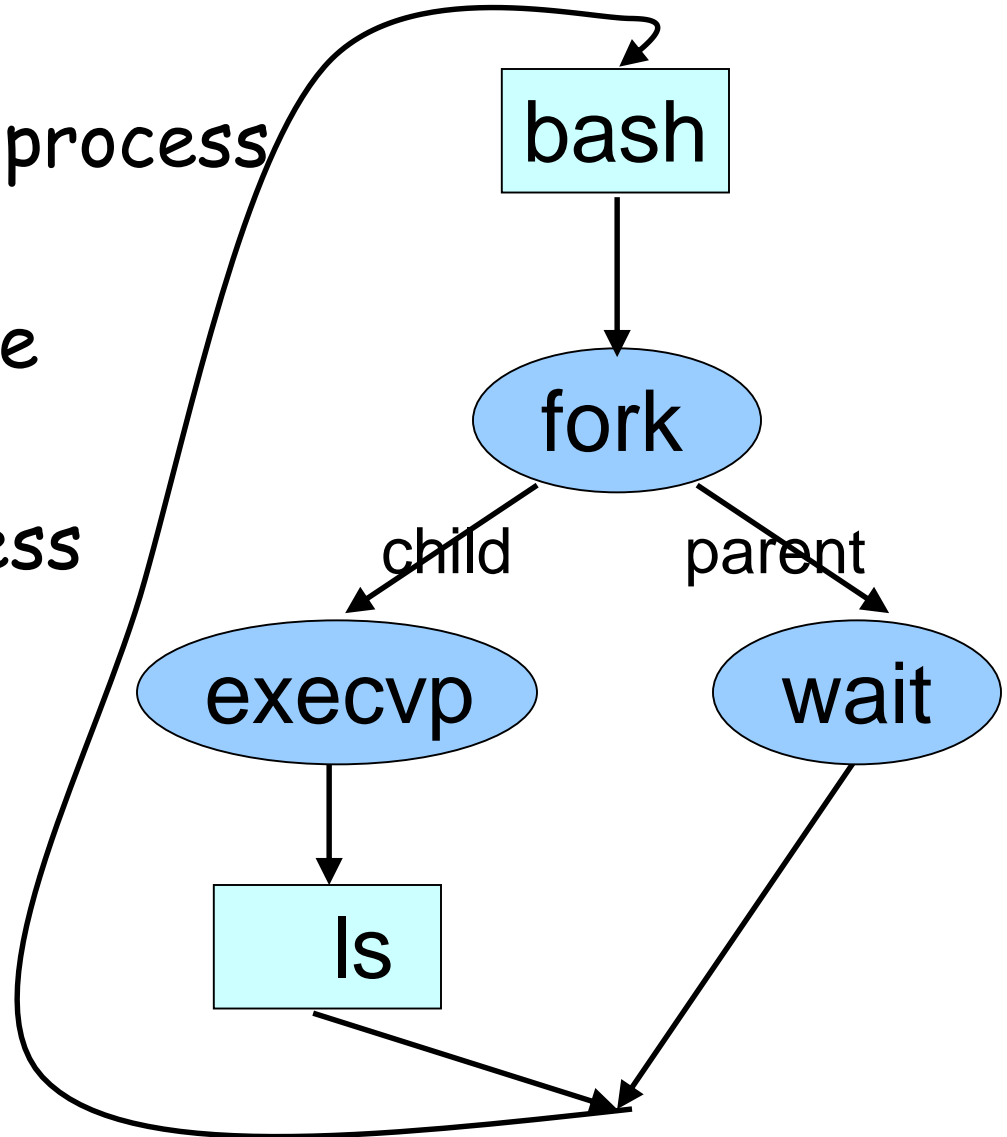
Waiting for the Child to Finish

- ❑ Parent should wait for children to finish.
 - Example: A shell waiting for operations to complete.
- ❑ Waiting for a child to terminate: `wait()`
 - Blocks until some child terminates.
 - Returns the process ID of the child process.
 - Or returns -1 if no children exist (i.e., already exited).
- ❑ Waiting for specific child to terminate: `waitpid()`
 - Blocks till a child with particular process ID terminates.

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);
```

Example: A Simple Shell

- ❑ Shell is the parent process
 - E.g., bash
- ❑ Parses command line
 - E.g., "ls -l"
- ❑ Invokes child process
 - `fork()`, `execvp()`
- ❑ Waits for child
 - `wait()`



Process Creation Using Fork

```
int main ()
{
    pid_t pid;
    int status = 0;
    pid = fork();
    if (pid < 0)
        perror("fork()");
    if (pid > 0) {
        /* parent */
        printf("I am parent\n");
    } else {
        /* child */
        printf("I am child\n");
        exit(status);
    }
}
```

The **fork** system call returns twice: it returns a zero to the child and the child process ID (pid) to the parent.

The **perror** function produces a message on the standard error output describing the last error encountered during a call to a system or library function (man page).

pid is zero which indicates a child process.

Fork System Call

- ❑ If `fork()` succeeds it returns the child PID to the parent and returns 0 to the child.
- ❑ If `fork()` fails, it returns -1 to the parent (no child is created) and sets `errno`.
- ❑ A program almost always uses this difference to do different things in the parent and child processes.
- ❑ Failure occurs when the limit of processes that can be created is reached.
- ❑ `pid_t` is defined as `typedef int pid_t`
 - We will use `pid_t` and `int` interchangeably.
- ❑ Other calls:
 - `int getpid()` - returns the PID of calling process.
 - `int getppid()` - returns the PID of parent process.

Wait

- ❑ Parents waits for a child (system call).
 - Blocks until a child terminates.
 - Returns pid of the child process.
 - Returns -1 if no child process exists (already exited).
 - status

`#include <sys/types.h>`

`#include <sys/wait.h>`

`pid_t wait(int *status)`

- ❑ Parent waits for a specific child to terminate.

`pid_t waitpid(pid_t pid, int *status, int options)`

Process Creation Using Fork

```
int main ()
{
    pid_t pid;
    int status = 0;
    pid = fork();
    if (pid < 0)
        perror("fork()");
    if (pid > 0) {
        /* parent */
        printf("I am parent\n");
        pid = wait(&status);
    } else {
        /* child */
        printf("I am child\n");
        exit(status);
    }
}
```

The **fork** syscall returns twice: it returns a zero to the child and the child process ID (pid) to the parent.

Parent uses **wait** to sleep until the child exits; **wait** returns child pid and status.

wait variants allow wait on a specific child, or notification of stops and other signals.

fork() Example

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    int i;
    pid = fork();
    if( pid > 0 )
    {
        /* parent */
        for( i=0; i < 1000; i++ )
            printf("\t\tPARENT %d\n", i);
    }
    else
    {
        /* child */
        for( i=0; i < 1000; i++ ) {
            printf( "CHILD %d\n", i );
        }
        return 0;
    }
}
```

What is the possible output?

fork () Example: Possible Output

PARENT 0

PARENT 1

PARENT 2

PARENT 3

PARENT 4

PARENT 5

PARENT 6

PARENT 7

PARENT 8

PARENT 9

CHILD 0

CHILD 1

CHILD 2

CHILD 3

CHILD 4

CHILD 5

CHILD 6

CHILD 7

CHILD 8

CHILD 9

fork () Example: Possible Output

PARENT 0
PARENT 1
PARENT 2
PARENT 3
PARENT 4
PARENT 5
PARENT 6

CHILD 0
CHILD 1
CHILD 2

PARENT 7
PARENT 8
PARENT 9

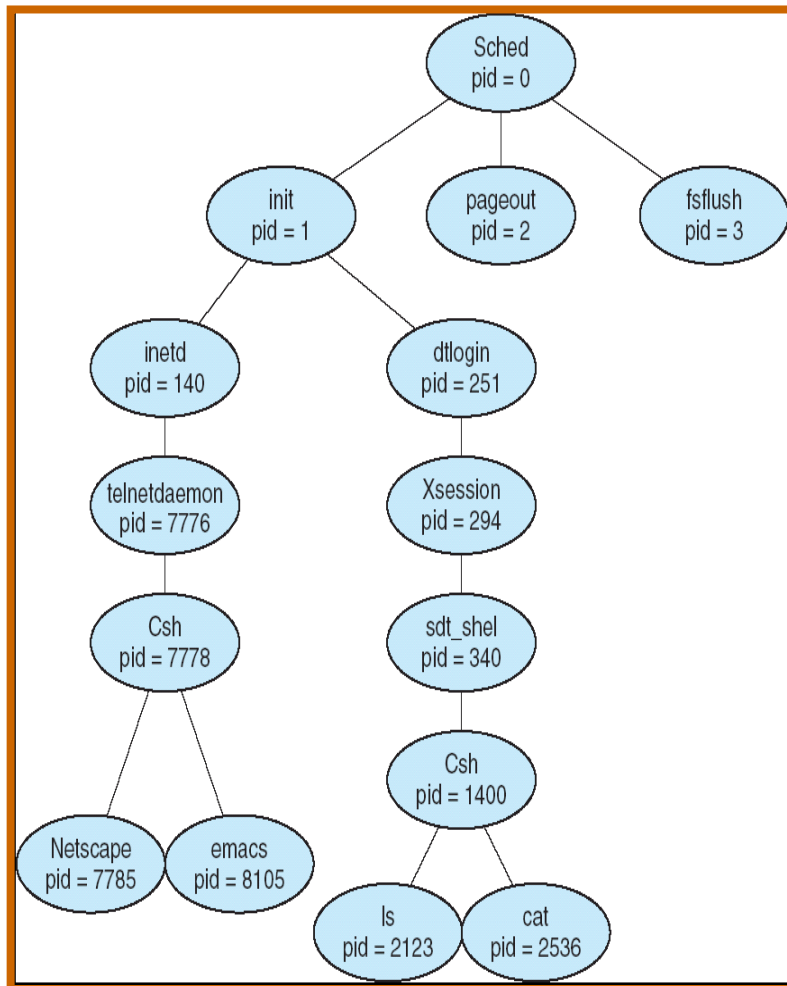
CHILD 3
CHILD 4
CHILD 5
CHILD 6
CHILD 7
CHILD 8
CHILD 9

Lots of possible outputs!!

Execution

- ❑ Processes get a share of the CPU before giving it up to give another process a turn.
- ❑ The switching between the parent and child depends on many factors:
 - machine load, system process scheduling.
- ❑ Output interleaving is **nondeterministic**
 - Cannot determine output by looking at code.

More about Process Operations



- ❑ In Unix-based systems, a hierarchy of processes is formed.
- ❑ In Unix, we can obtain a listing of processes by using the **ps** command.
- ❑ **ps -el** will list complete information for all processes.

Exec

- ❑ The term **exec** refers to a family of functions where each of the functions replace a process's program (the one calling one of the exec functions) with a new loaded program.
- ❑ A call to a function from **exec** loads a binary file into memory (destroying the memory image of the program calling it).
- ❑ The new program starts executing from the beginning (where main begins).
- ❑ On success, **exec** never returns; on failure, **exec** returns -1.
- ❑ The different versions are different primarily in the way parameters are passed.

Exec

- ❑ The exec family consists of these functions: `execvp`, `execlp`, `execv`, `execve`, `execl`, `execle`.
 - Functions with *p* in their name (`execvp`, `execlp`) search for the program in current path; functions without *p* must be given full path.
 - Functions with *v* in their name (`execv`, `execvp`, `execve`) differ from functions with "l" (`execl`, `execlp`, `execle`) in the way arguments are passed.
 - Functions with *e* accept array of environment variables.

Versions of exec

- Versions of exec offered by C library:

```
int execl( const char *path, const char *arg, ... );
int execlp( const char *file, const char *arg, ... );
int execl( const char *path, const char *arg
           , ..., char *const envp[] );
int execv( const char *path, char *const argv[] );
int execvp( const char *file, char *const argv[] );
int execve( const char *filename, char *const argv [], char *const
           envp[] );
```

Exec Example

Program A:

```
int i = 5;
printf("%d\n",i);
execl("B", "B", NULL);
printf("%d\n",i);
```

Program B:

```
main()
{
    printf("hello\n");
}
```

- What is the output of program A?

5
hello

- Why is it not this?

5
hello
5

- The **exec** command replaces the instructions in the process with instructions for program B. It starts at the first instruction (starts at **main**).

Exec Example

Program A:

```
int i = 5;
prog_argv[0] = "B";
prog_argv[1] = NULL;
printf("%d\n",i);
execv(prog_argv[0],
      prog_arg);
printf("%d\n",i);
```

Program B:

```
main()
{
    printf("hello\n");
}
```

- ❑ Same functionality as the program on the previous slide.
- ❑ Used **execv** instead of **execl**.
- ❑ **execv** uses an array to pass arguments.
- ❑ **execl** uses a list to pass arguments.
- ❑ If you use **execv** you must provide the full path name e.g.,

```
prog_argv[0] =
    "/eecs/courses/OS/CodeE
    xamples/B"
```

Exec Example

```
int main(int argc, char *argv[])
{

    char *prog1_argv[4];
    int i = 5;

    execlp("ls","ls","-l","a.c",NULL);

    perror("execlp\n");
    printf("%d\n",i);

}
```

- ❑ In this example, note that the command is
`ls -l a.c`
- ❑ Each argument is in the list.
- ❑ Question:
 - What would cause the **perror** function to be executed?

Exec Example

```
int main(int argc, char *argv[])
{

    char *prog1_argv[4];
    int i = 5;

    prog1_argv[0] = "ls";
    prog1_argv[1] = "-l";
    prog1_argv[2] = "a.c";
    prog1_argv[3] = NULL;

    execvp(prog1_argv[0],
           prog1_argv);

    perror("execvp\n");
    printf("%d\n", i);

}
```

- Same example as that on the previous slide but **execvp** is used which requires an array.

Fork and Exec

- ❑ Child process may choose to execute some other program than the parent by using one of the exec calls.
- ❑ `exec` overlays a new program on the existing process.
- ❑ Child will not return to the old program unless `exec` fails. This is an important point to remember.
- ❑ File descriptors are preserved.

Example

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    pid = fork();
    if (pid < 0)
        perror("fork()");
```

```
    if (pid > 0)
    {
        wait(NULL);
        printf("Child Complete");
    } else{
        if (pid == 0)
            execlp("ls","ls", "-l", "a.c",
                  NULL);
    }
}
```