

Deadlocks

Overview


- ❑ Resources.
- ❑ Why do deadlocks occur?
- ❑ Dealing with deadlocks.
 - Ignoring them: ostrich algorithm.
 - Detecting.

The Deadlock Problem

- ❑ A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set
- ❑ Example
 - System has 2 disk drives
 - P_1 and P_2 each hold one disk drive and each needs another one
- ❑ Example
 - semaphores A and B , initialized to 1

P_0	P_1
wait (A);	wait(B)
wait (B);	wait(A)


Bridge Crossing Example




- ❑ Traffic only in one direction
- ❑ Each section of a bridge can be viewed as a resource
- ❑ If a deadlock occurs, it can be resolved if one car backs up
 - preempt resources and rollback
- ❑ Several cars may have to be backed up if a deadlock occurs
- ❑ Starvation is possible
- ❑ **Note** : Most OSes do not prevent or deal with deadlocks

Deadlock Principles

- A deadlock is a permanent blocking of a set of threads
✓ a deadlock can happen while threads/processes are competing for system resources or communicating with each other



(a) Deadlock possible



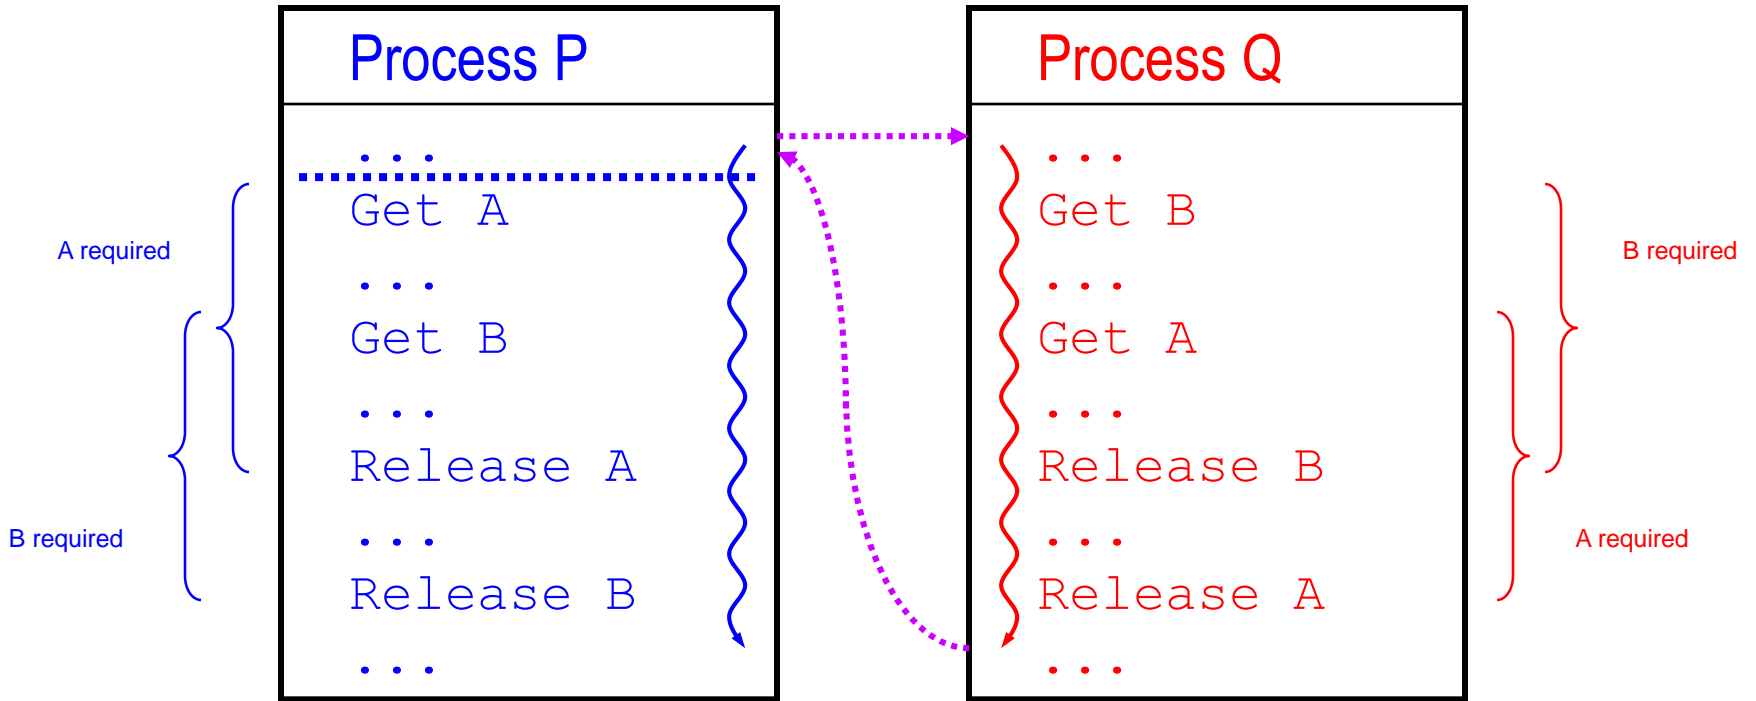
(b) Deadlock

Illustration of a deadlock

Deadlock Principles

➤ Illustration of a deadlock – scheduling path 1 ☺

- ✓ Q executes everything before P can ever get A
- ✓ when P is ready, resources A and B are free and P can proceed

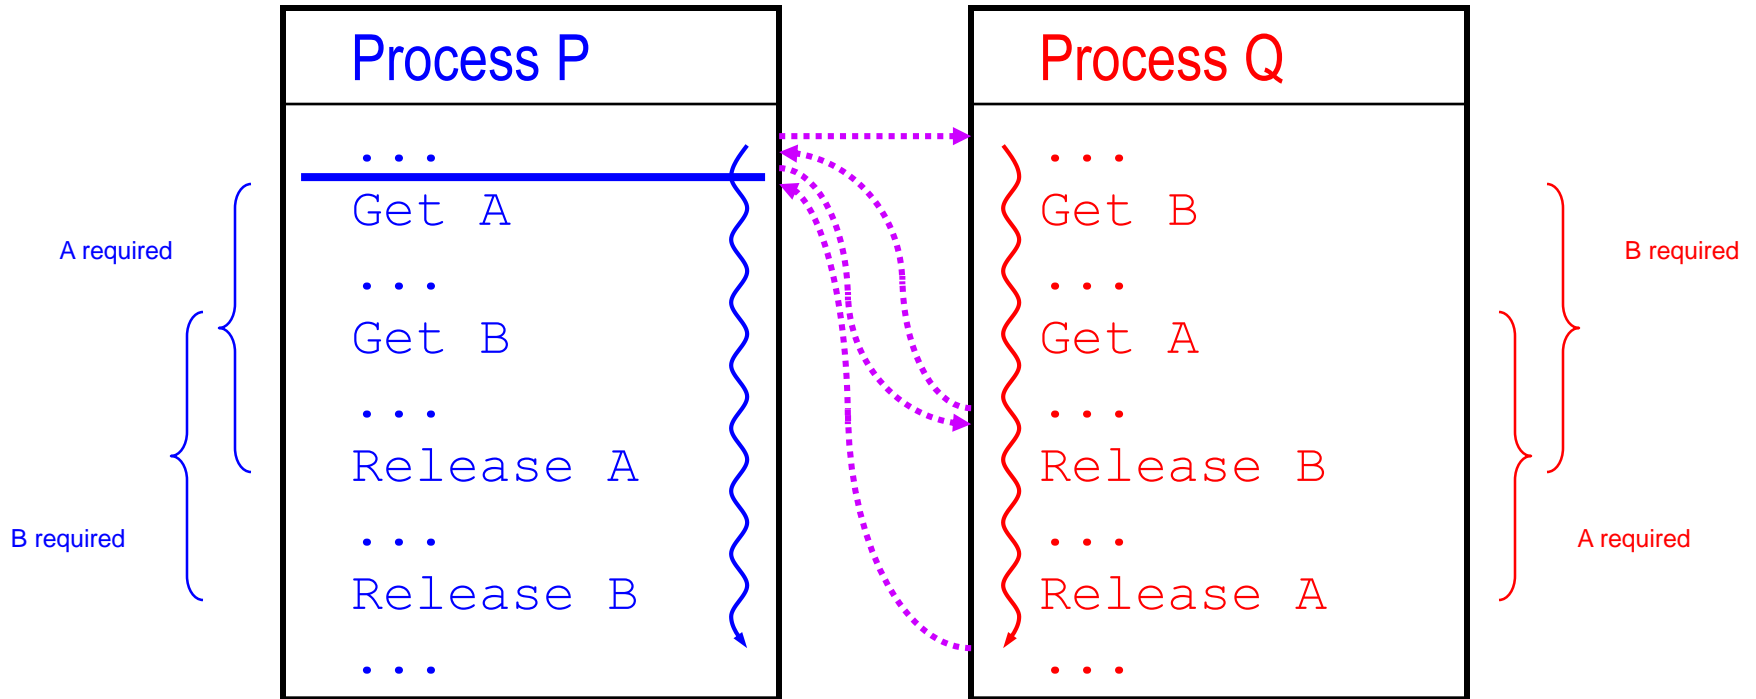


Happy scheduling 1

Deadlock Principles

➤ Illustration of a deadlock — scheduling path 2 ☺

✓ Q gets B and A, then P is scheduled; P wants A but is blocked by A's mutex; so Q resumes and releases B and A; P can now go

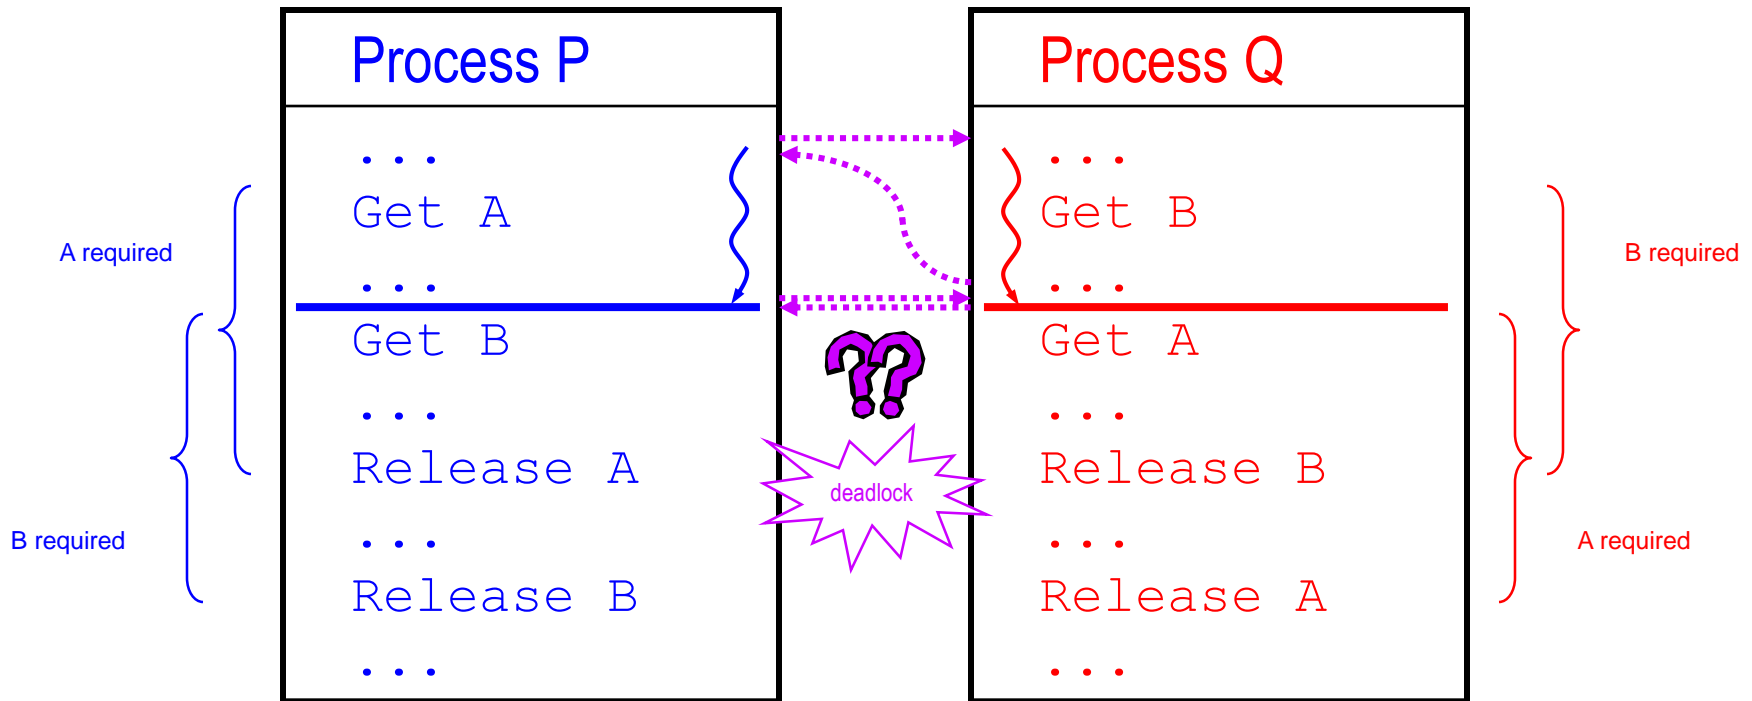


Happy scheduling 2

Deadlock Principles

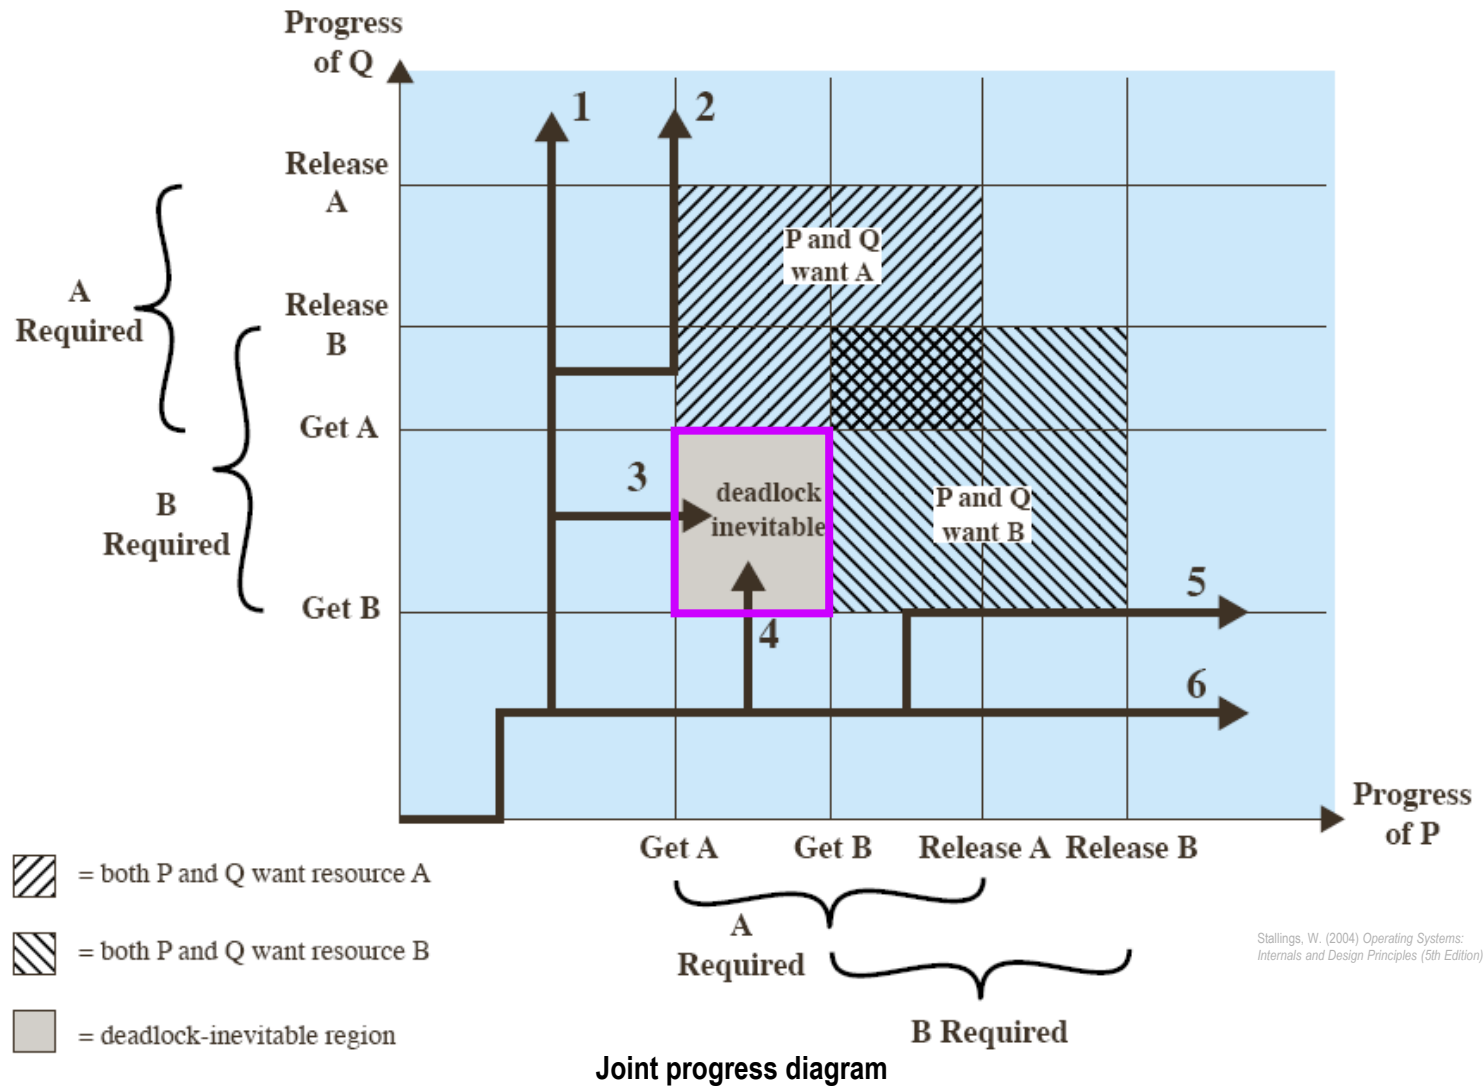
Illustration of a deadlock — scheduling path 3 ☹️

Q gets only B, then P is scheduled and gets A; now both P and Q are blocked, each waiting for the other to release a resource



Bad scheduling → **deadlock**

Deadlock Principles




Deadlock Principles

➤ Deadlocks depend on the program and the scheduling


✓ **program design**

- the order of the statements in the code creates the "landscape" of the joint progress diagram
- this landscape may contain gray "swamp" areas leading to **deadlock**



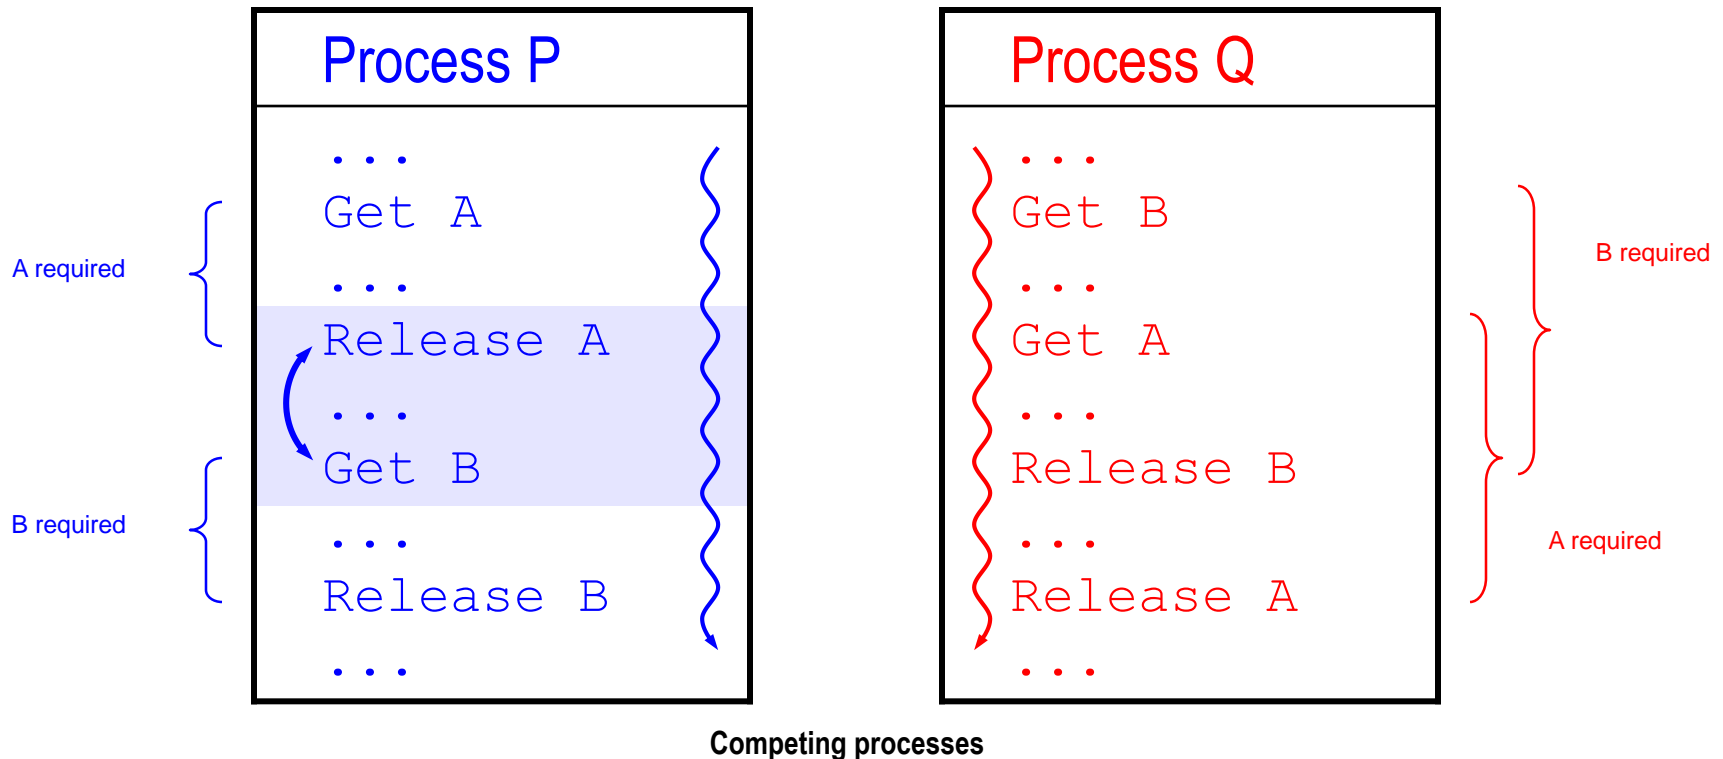
✓ **scheduling condition**

- the interleaved dynamics of multiple executions traces a "path" in this landscape
- this path may sink in the swamps

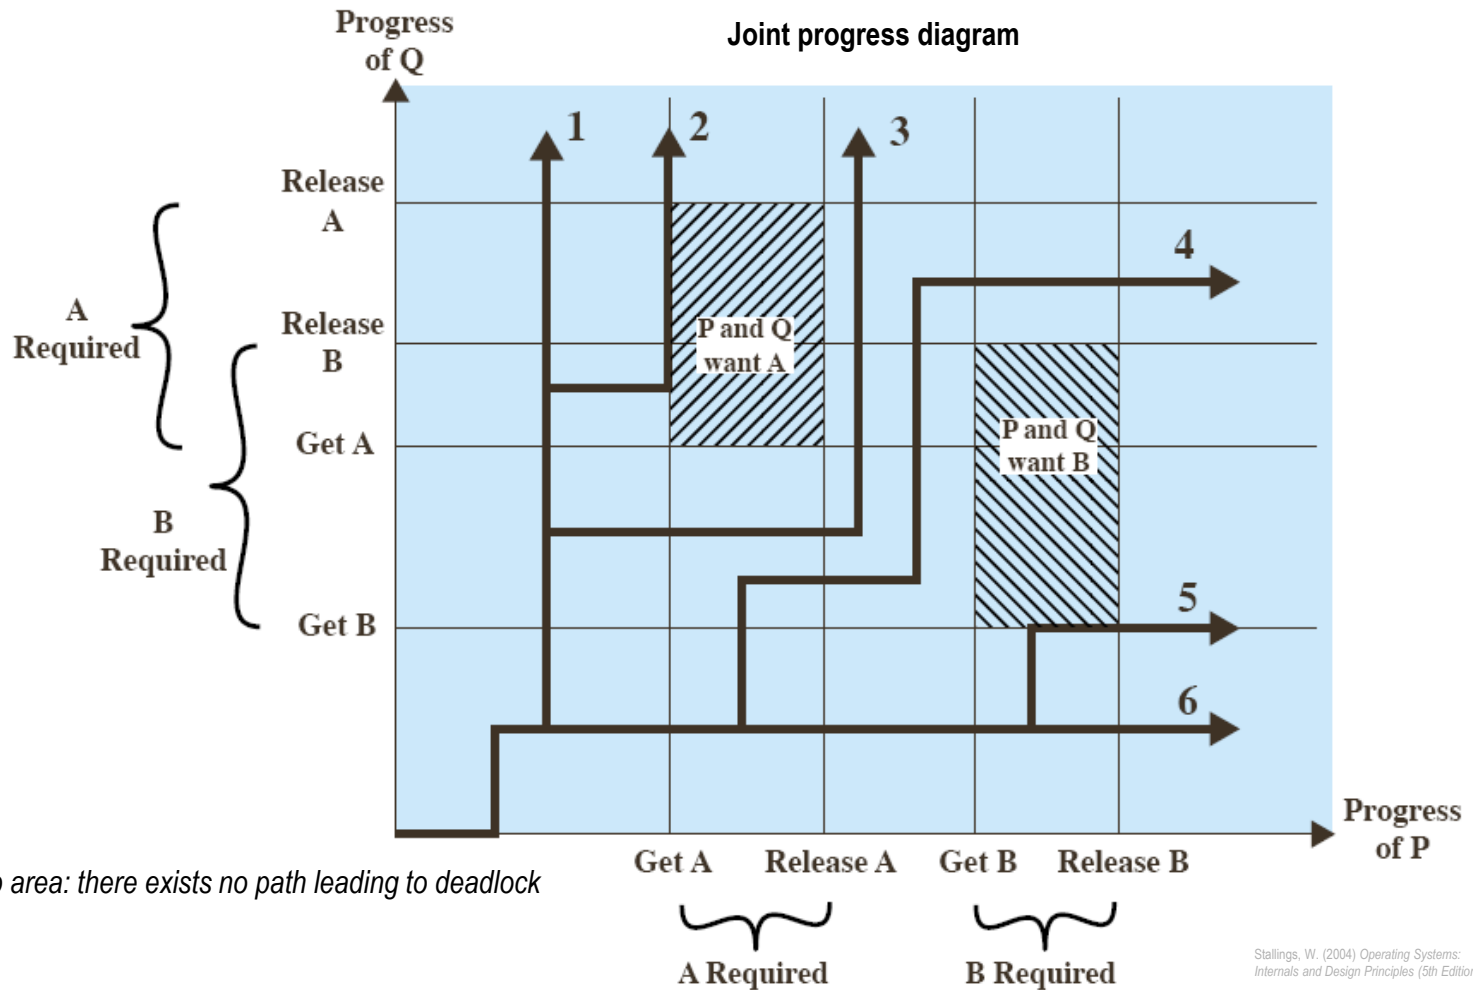


Deadlock Principles

- Changing the program changes the landscape
- ✓ here, P releases A before getting B
 - ✓ deadlocks between P and Q are not possible anymore



Deadlock Principles



System Model

- ❑ Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- ❑ Each resource type R_i has W_i instances.
- ❑ Each process utilizes a resource as follows:
 - request
 - use
 - release

Four conditions for deadlock

❑ Mutual exclusion

- Each resource is assigned to at most one process

❑ Hold and wait

- A process holding resources can request more resources

❑ No preemption

- Previously granted resources cannot be forcibly taken away

❑ Circular wait

- There must be a circular chain of 2 or more processes where each is waiting for a resource held by the next member of the chain

Deadlock Characterization


- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Resource-Allocation Graph


- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- **request edge** - directed edge $P_i \rightarrow R_j$
- **assignment edge** - directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (Cont.)


□ Process




□ Resource Type with 4 instances




□ P_i requests instance of R_j




□ P_i is holding an instance of R_j




Example of a Resource Allocation Graph



Resource Allocation Graph With A Deadlock



Graph With A Cycle But No Deadlock



Basic Facts

- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, possibility of deadlock

Resources

- ❑ Resource: something a process uses
 - Usually limited (at least somewhat)
- ❑ Examples of computer resources
 - Printers
 - Semaphores / locks
 - Tables (in a database)
- ❑ Processes need access to resources in reasonable order
- ❑ Two types of resources:
 - Preemptable resources: can be taken away from a process with no ill effects
 - Nonpreemptable resources: will cause the process to fail if taken away


When do deadlocks happen?

□ Suppose

- Process 1 holds resource A and requests resource B
- Process 2 holds B and requests A
- Both can be blocked, with neither able to proceed

□ Deadlocks occur when ...

- Processes are granted exclusive access to devices or software constructs (resources)
- Each deadlocked process needs a resource held by another deadlocked process



DEADLOCK!

Using resources

- ❑ Sequence of events required to use a resource
 - Request the resource
 - Use the resource
 - Release the resource
- ❑ Can't use the resource if request is denied
 - Requesting process has options
 - Block and wait for resource
 - Continue (if possible) without it: may be able to use an alternate resource
 - Process fails with error code
 - Some of these may be able to prevent deadlock...

What is a deadlock?

- Formal definition:




“A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.”

- Usually, the event is release of a currently held resource

- In deadlock, none of the processes can

- Run
- Release resources
- Be awakened

Resource allocation graphs



- ❑ Resource allocation modeled by directed graphs

- ❑ Example 1:

- Resource R assigned to process A

- ❑ Example 2:

- Process B is requesting / waiting for resource S

- ❑ Example 3:

- Process C holds T, waiting for U
- Process D holds U, waiting for T
- C and D are in deadlock!


Dealing with deadlock

- ❑ How can the OS deal with deadlock?
 - Ignore the problem altogether!
 - Hopefully, it'll never happen...
 - Detect deadlock & recover from it
 - Dynamically avoid deadlock
 - Careful resource allocation
 - Prevent deadlock
 - Remove at least one of the four necessary conditions
- ❑ We'll explore these tradeoffs


Getting into deadlock

A

Acquire R
Acquire S
Release R
Release S




Acquire R




Acquire S

B

Acquire S
Acquire T
Release S
Release T




Acquire S




Acquire T

C

Acquire T
Acquire R
Release T
Release R



Acquire T



Acquire R

Deadlock!

Not getting into deadlock...

- ❑ Many situations *may* result in deadlock (but don't have to)
 - In previous example, A could release R before C requests R, resulting in no deadlock
 - Can we always get out of it this way?
- ❑ Find ways to:
 - Detect deadlock and reverse it
 - Stop it from happening in the first place


The Ostrich Algorithm

- ❑ Pretend there's no problem
- ❑ Reasonable if
 - Deadlocks occur very rarely
 - Cost of prevention is high
- ❑ UNIX and Windows take this approach
 - Resources (memory, CPU, disk space) are plentiful
 - Deadlocks over such resources rarely occur
 - Deadlocks typically handled by rebooting
- ❑ Trade off between convenience and correctness

Detecting deadlocks using graphs

- ❑ Process holdings and requests in the table and in the graph (they're equivalent)
- ❑ Graph contains a cycle \Rightarrow deadlock!
 - Easy to pick out by looking at it (in this case)
 - Need to mechanically detect deadlock
- ❑ Not all processes are deadlocked (A, C, F not in deadlock)

Process	Holds	Wants
A	R	S
B		T
C		S
D	U	S,T
E	T	V
F	W	S
G	V	U



Deadlock detection algorithm

- ❑ General idea: try to find cycles in the resource allocation graph
- ❑ Algorithm: depth-first search at each node
 - Mark arcs as they're traversed
 - Build list of visited nodes
 - If node to be added is already on the list, a cycle exists!
- ❑ Cycle == deadlock

```
For each node N in the graph {  
  Set L = empty list  
  unmark all arcs  
  Traverse (N,L)  
}  
If no deadlock reported by now, there isn't any  
  
define Traverse (C,L) {  
  If C in L, report deadlock!  
  Add C to L  
  For each unmarked arc from C {  
    Mark the arc  
    Set A = arc destination  
    /* NOTE: L is a  
       local variable */  
    Traverse (A,L)  
  }  
}
```


Deadlock Avoidance


- ❑ makes sure the system stays in a safe state if a request is satisfied
- ❑ prevents circular waits
- ❑ requires additional information at the beginning of the execution of a process
 - maximum number of instances per resource type for each process

Safe State

- A system is in a **safe state** when given the current allocation the system is able to handle any number of requests, in some order, without getting into a deadlock.

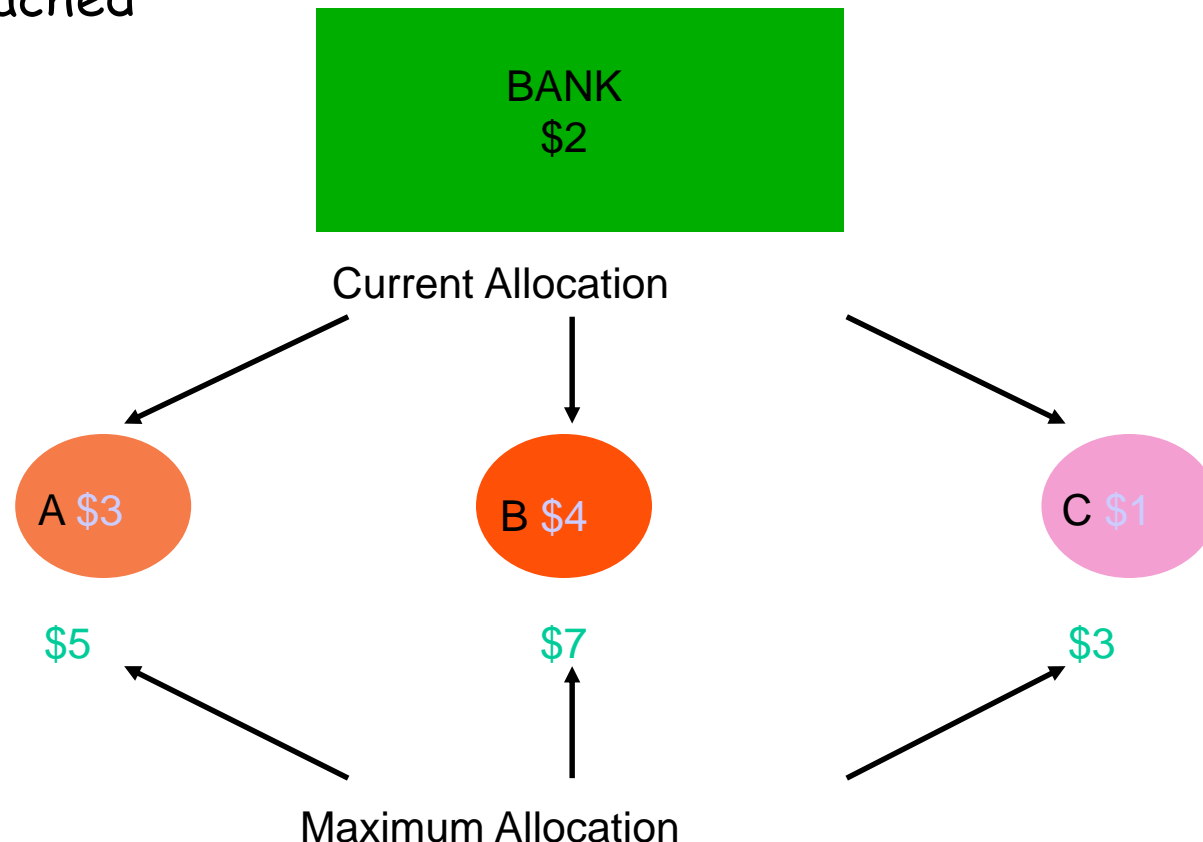
Example

- ❑ Bank gives loans to customers
 - maximum allocation = credit limit



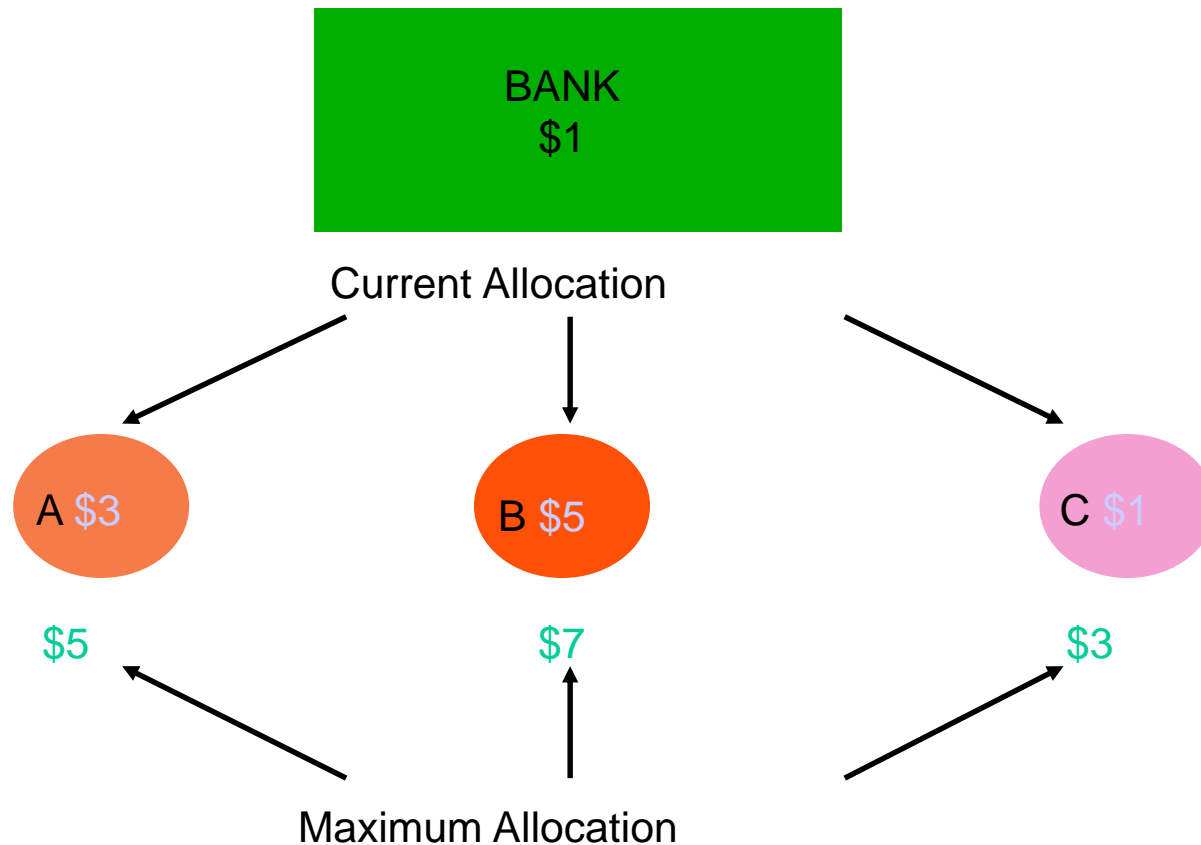
❑ Safe State?

- Will the bank be able to give each customer a loan up to the full credit limit?
 - not necessarily all customers simultaneously
 - order is not important
 - customers will pay back their loan once their credit limit is reached



❑ Still Safe?

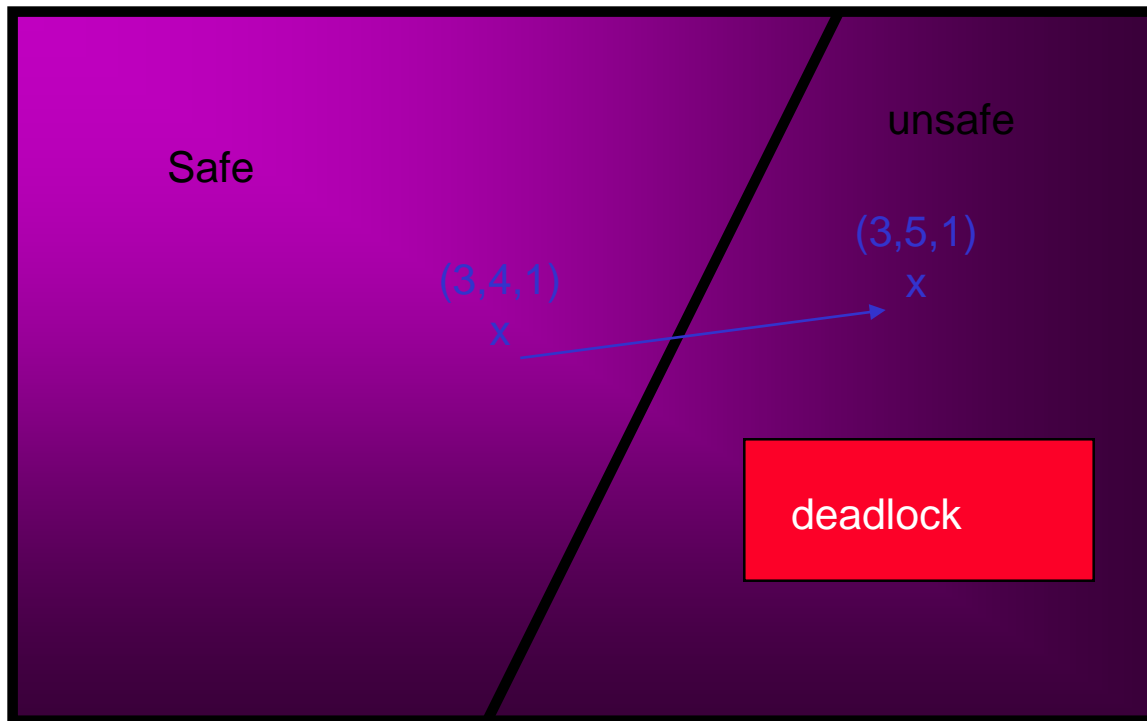
- after customer B requests and is granted \$1, is the bank still safe?



Safe State Space



Bank Safe State Space



Banker's Algorithm

- ❑ before a request is granted, check the system's state
 - assume the request is granted
 - if it is still safe, the request can be honored
 - otherwise the process has to wait
 - overly careful
 - there are cases when the system is unsafe, but not in a deadlock

Example Banker's Algorithm

- Initially given: current allocation, maximum allocation, available resources

Process	Current Allocation	Maximum Allocation	Current Work Available	Remaining Needed = Max - Aval	Total Available
	A B C	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3			10 5 7
P1	2 0 0	3 2 2			
P2	3 0 2	9 0 2			
P3	2 1 1	4 2 2			
P4	0 0 2	5 3 3			
	7 2 5				

Example Banker's Algorithm

- ❑ First check which process is in the safe state.
- ❑ $Need_i \leq Work \Rightarrow Work = Work + Allocation$
 - If(RemainingNeeded <= CurrentWorkAvailable)
 $CurrentWorkAvailable = CurrentWorkAvailable + CurrentAllocation$
- ❑ Update the safe sequence.

Process	Current Allocation	Maximum Allocation	Current Work Available	Remaining Needed = Max - Aval	Total Available
	A B C	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3			10 5 7
P1	2 0 0	3 2 2			
P2	3 0 2	9 0 2			
P3	2 1 1	4 2 2			
P4	0 0 2	5 3 3			
	7 2 5				

Example Banker's Algorithm

□ Initially given current allocation, maximum allocation, available resources

○ If(RemainingNeeded <= CurrentWorkAvailable)

CurrentWorkAvailable = CurrentWorkAvailable + CurrentAllocation

Process	Current Allocation	Maximum Allocation	Current Work Available	Remaining Needed = Max - Aval	Total Available
	A B C	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3	3 3 2	7 4 3	10 5 7
P1	2 0 0	3 2 2	5 3 2	1 2 2	
P2	3 0 2	9 0 2	7 4 3	6 0 0	
P3	2 1 1	4 2 2	7 4 5	2 1 1	
P4	0 0 2	5 3 3	7 5 5	5 3 1	
	7 2 5		10 5 7		