

Basic Memory Management

Outline

- ❑ Memory management requirements
- ❑ Basic memory management
- ❑ Fixed Partitions
- ❑ Dynamic Partitions

In an ideal world...

- ❑ The ideal world has memory that is
 - Very large
 - Very fast
 - Non-volatile (doesn't go away when power is turned off)
- ❑ The real world has memory that is:
 - Very large
 - Very fast
 - Affordable!

⇒ Pick any two...
- ❑ Memory management goal: make the real world look as much like the ideal world as possible.

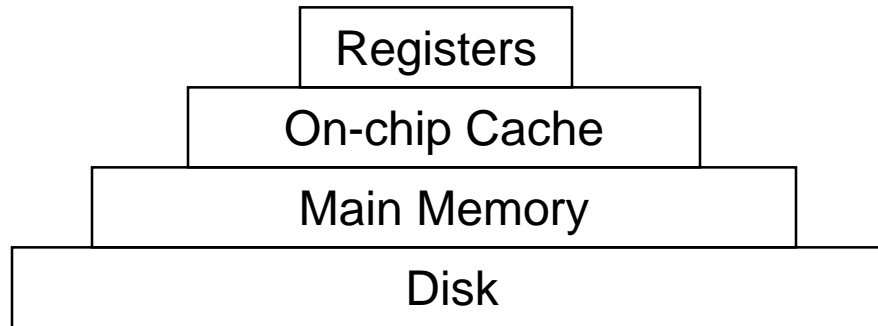
Introduction

- ❑ Our machines today have lots of more memory than the IBM 7094 - leading edge machine of the 1960's
- ❑ Cost of memory had dropped dramatically
- ❑ Bill Gates (former chair of Microsoft) once said "640K should be enough"

Introduction


- ❑ Software tends to expand to fill the memory available
 - The 1 terabyte of memory - the researchers already want more.
- ❑ Operating systems must manage memory

Memory Hierarchy

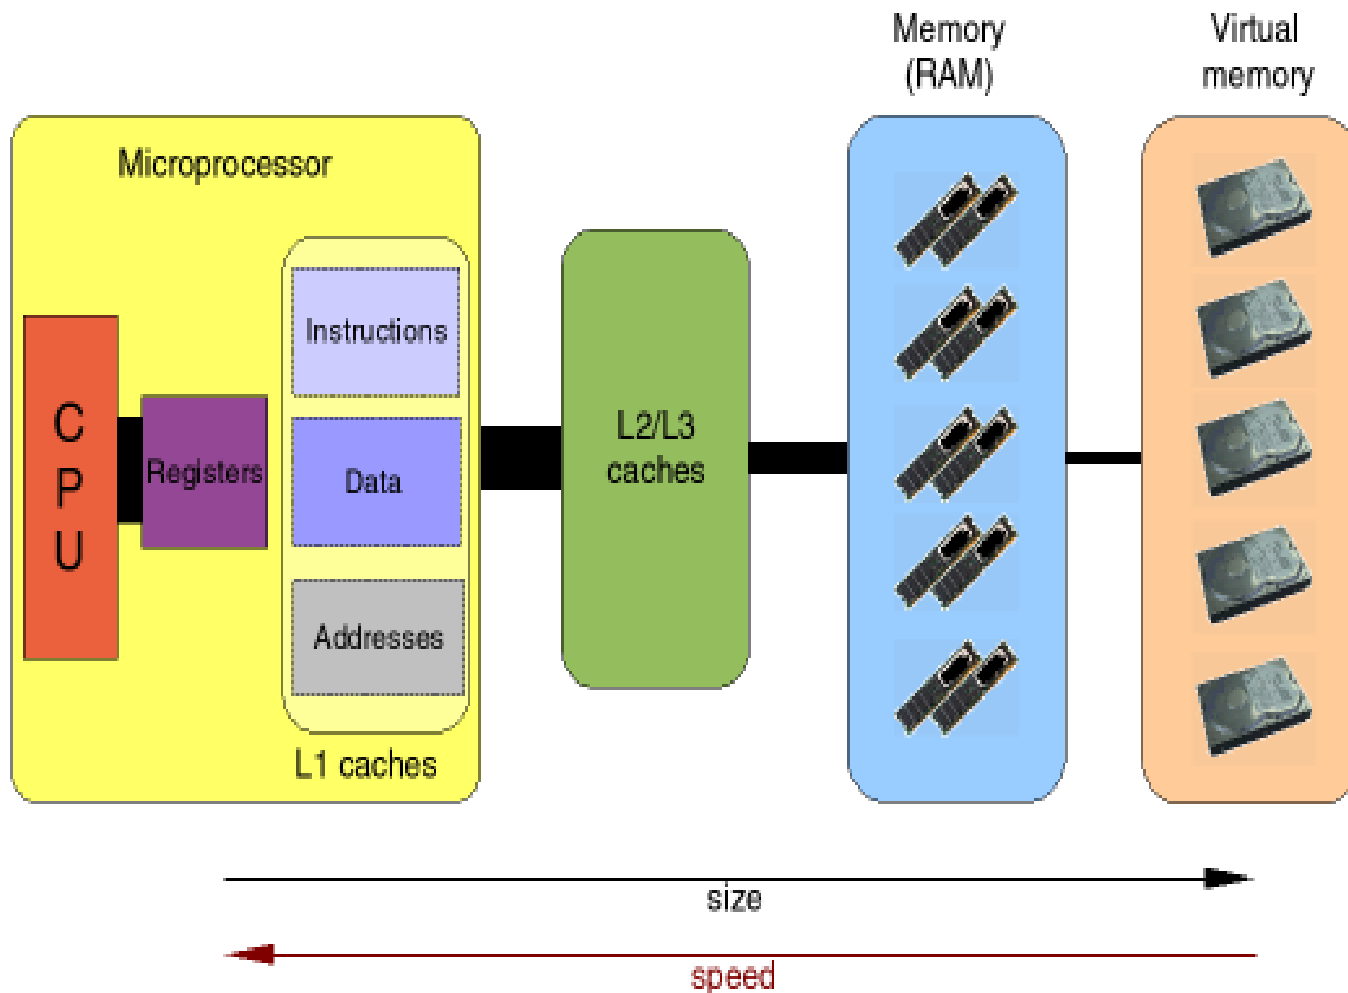


- ❑ A CPU waiting for data from main memory is not desired.
- ❑ Remedy: Add fast memory between the CPU and main memory called a **cache**.

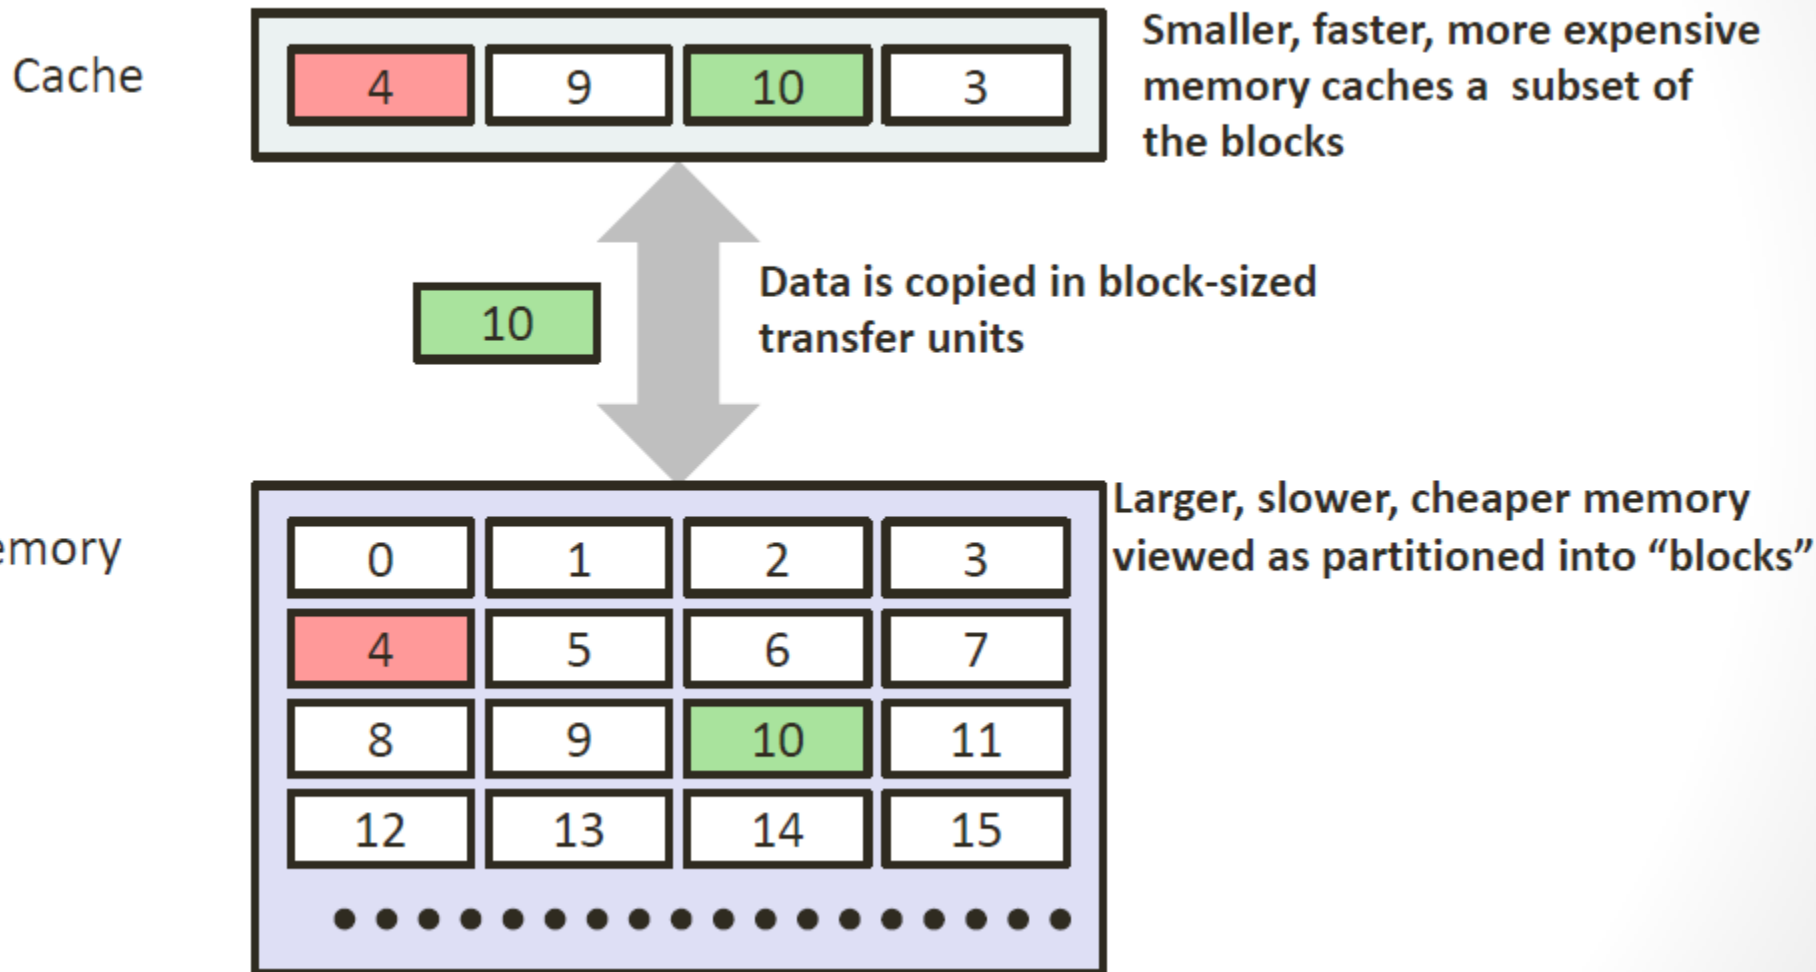
Memory Hierarchy



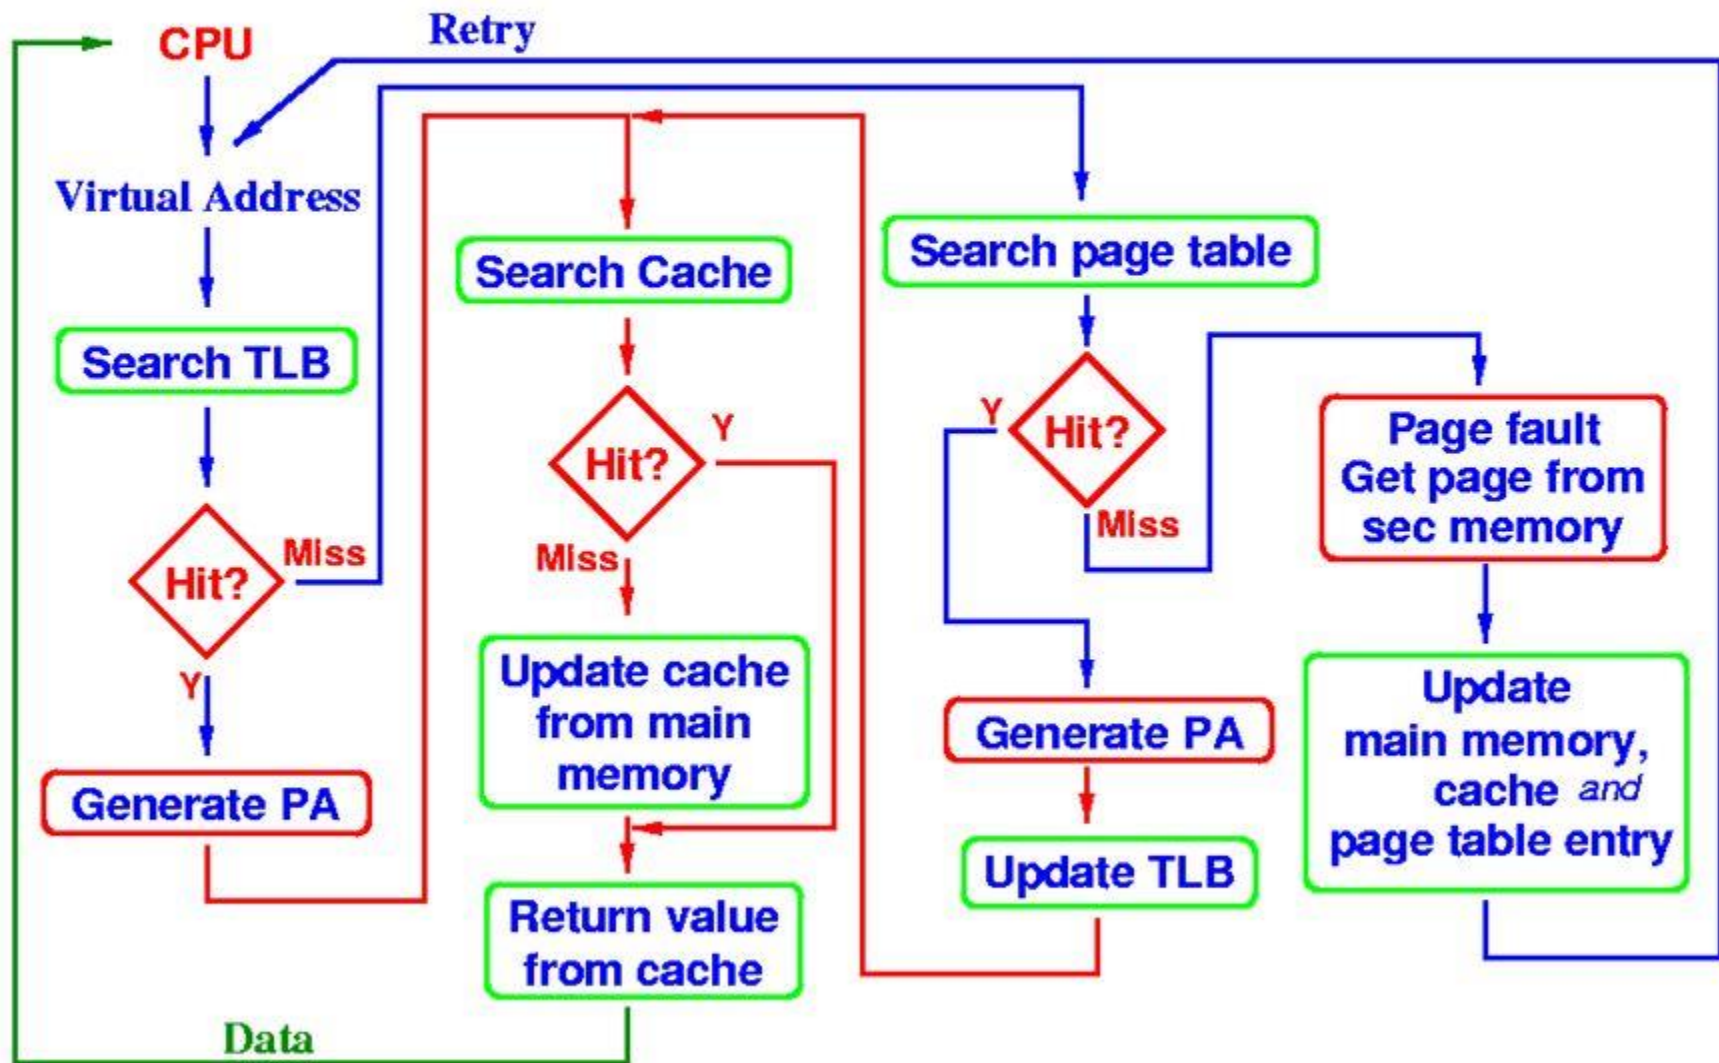
Memory Hierarchy



Memory Hierarchy



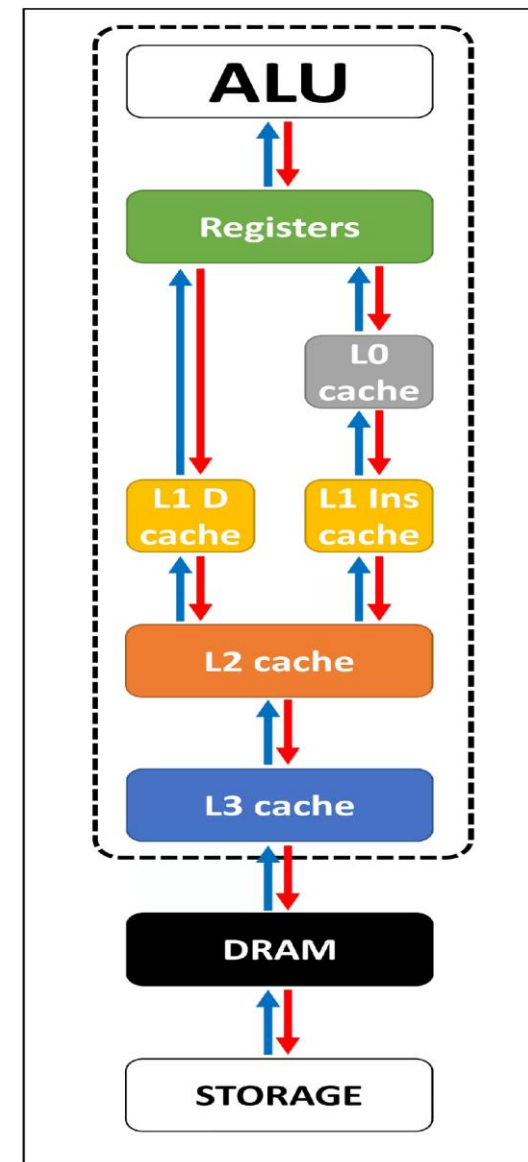
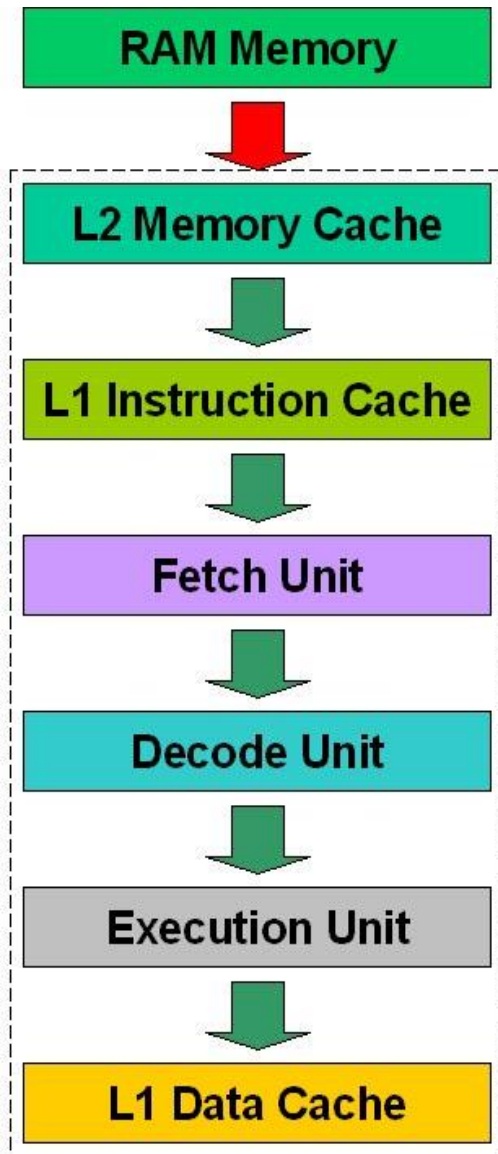
Memory Hierarchy - Operation



Memory hierarchy

- ❑ What is the memory hierarchy?
 - Different levels of memory
 - Some are small & fast
 - Others are large & slow
- ❑ What levels are usually included?
 - Cache: small amount of fast, expensive memory
 - L1 (level 1) cache: usually on the CPU chip
 - L2 & L3 cache: off-chip, made of SRAM in old machines
 - Main memory: medium-speed, medium price memory (DRAM)
 - Disk: many gigabytes of slow, cheap, non-volatile storage
- ❑ Memory manager handles the memory hierarchy

Memory hierarchy



Basic memory management

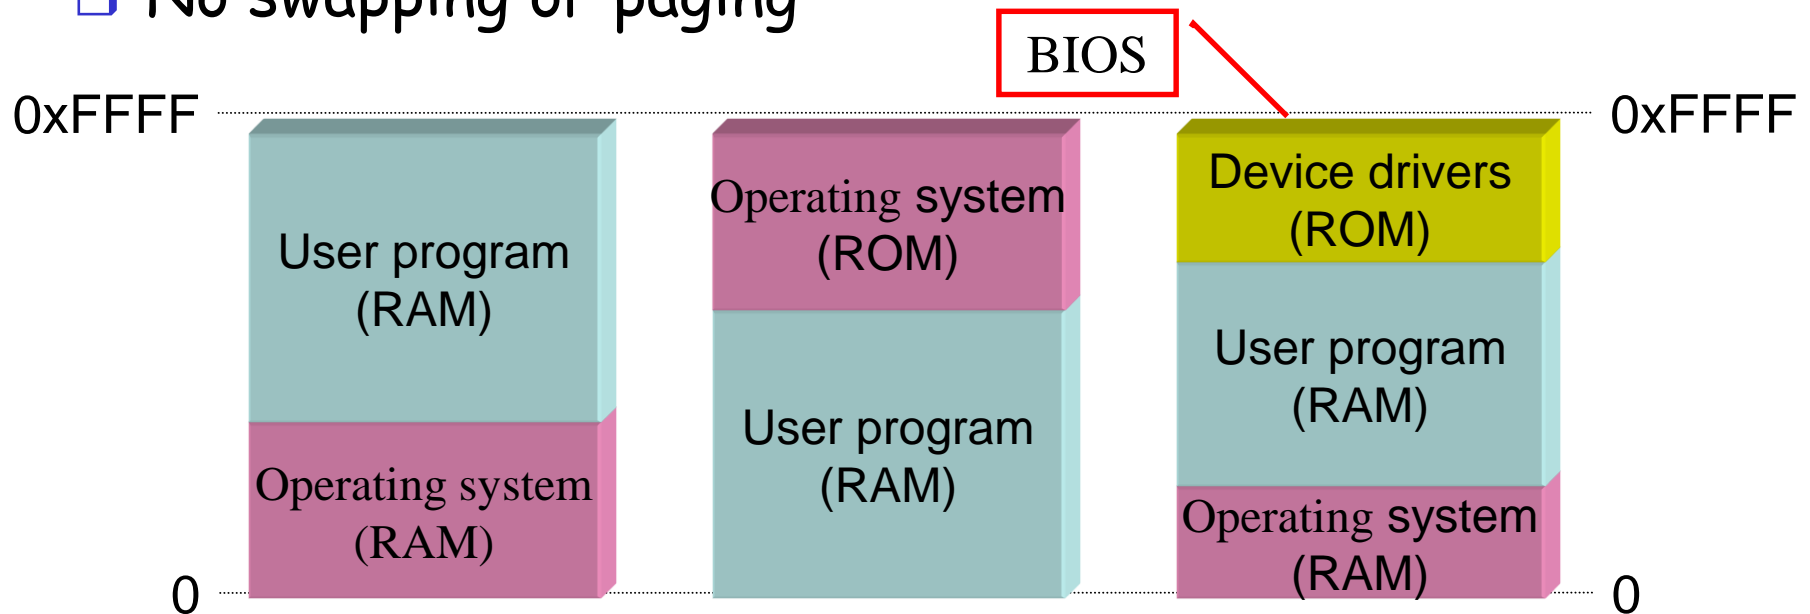
□ Components include

- Operating system (perhaps with device drivers)
- Single process

□ Goal: lay these out in memory


- Memory protection may not be an issue (only one program)
- Flexibility may still be useful (allow OS changes, etc.)

□ No swapping or paging



Fixed Partitioning


- ❑ Main memory use is inefficient.
- ❑ Any program, no matter how small, occupies an entire partition.
 - This is **internal fragmentation**.



Placement Algorithm with Partitions


- ❑ Equal-size partitions
 - Since all partitions are of equal size, it does not matter which partition is used
- ❑ Unequal-size partitions
 - Each process can be assigned to the smallest partition within which it will fit
 - There is a queue for each partition
 - Processes are assigned in such a way as to minimize wasted memory within a partition

One Process Queue per Partition



One Common Process Queue

- When its time to load a process into main memory the smallest available partition that will hold the process is selected




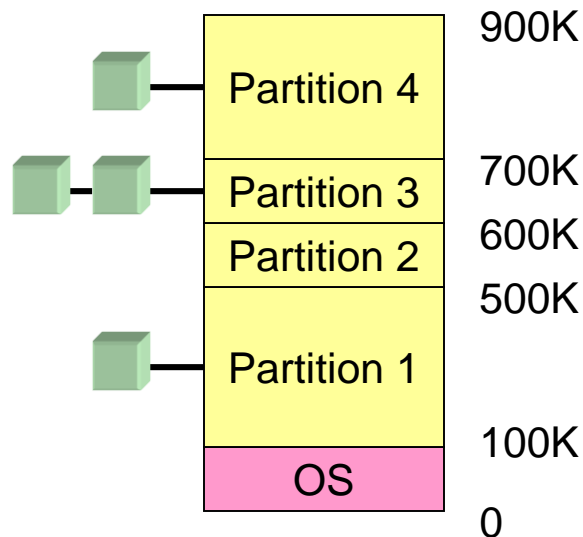
Fixed partitions: multiple programs

❑ Fixed memory partitions

- Divide memory into fixed spaces
- Assign a process to a space when it's free

❑ Mechanisms

- Separate input queues for each partition
- Single input queue: better ability to optimize CPU usage




How many programs is enough?

- ❑ Several memory partitions (fixed or variable size)
- ❑ Lots of processes wanting to use the CPU
- ❑ Tradeoff
 - More processes utilize the CPU better
 - Fewer processes use less memory (cheaper!)
- ❑ How many processes do we need to keep the CPU fully utilized?
 - This will help determine how much memory we need
 - Is this still relevant with memory costing \$100/16GB?

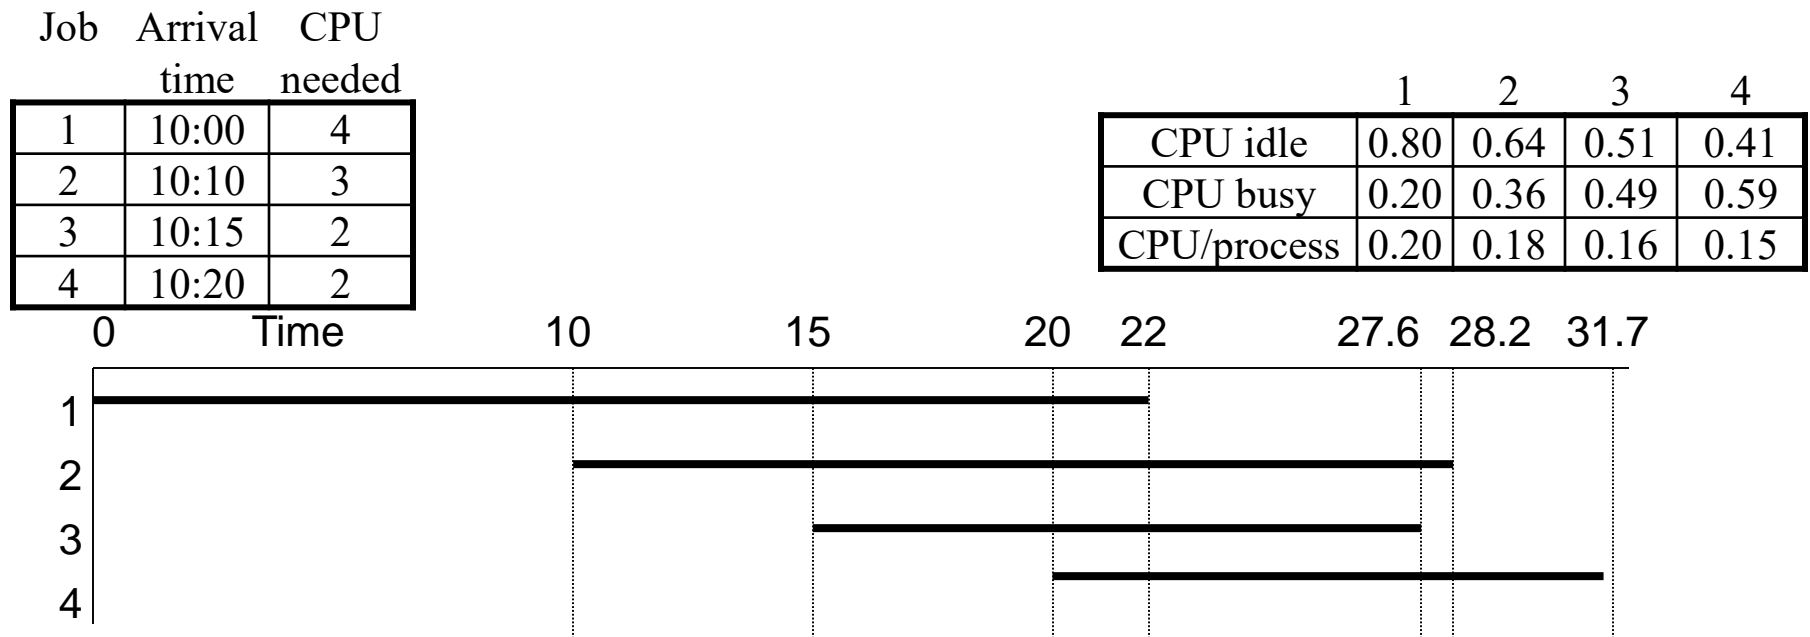
Modeling multiprogramming

- More I/O wait means less processor utilization
 - At 20% I/O wait, 3-4 processes fully utilize CPU
 - At 80% I/O wait, even 10 processes aren't enough
- This means that the OS should have more processes if they're I/O bound
- More processes => memory management & protection more important!



Multiprogrammed system performance

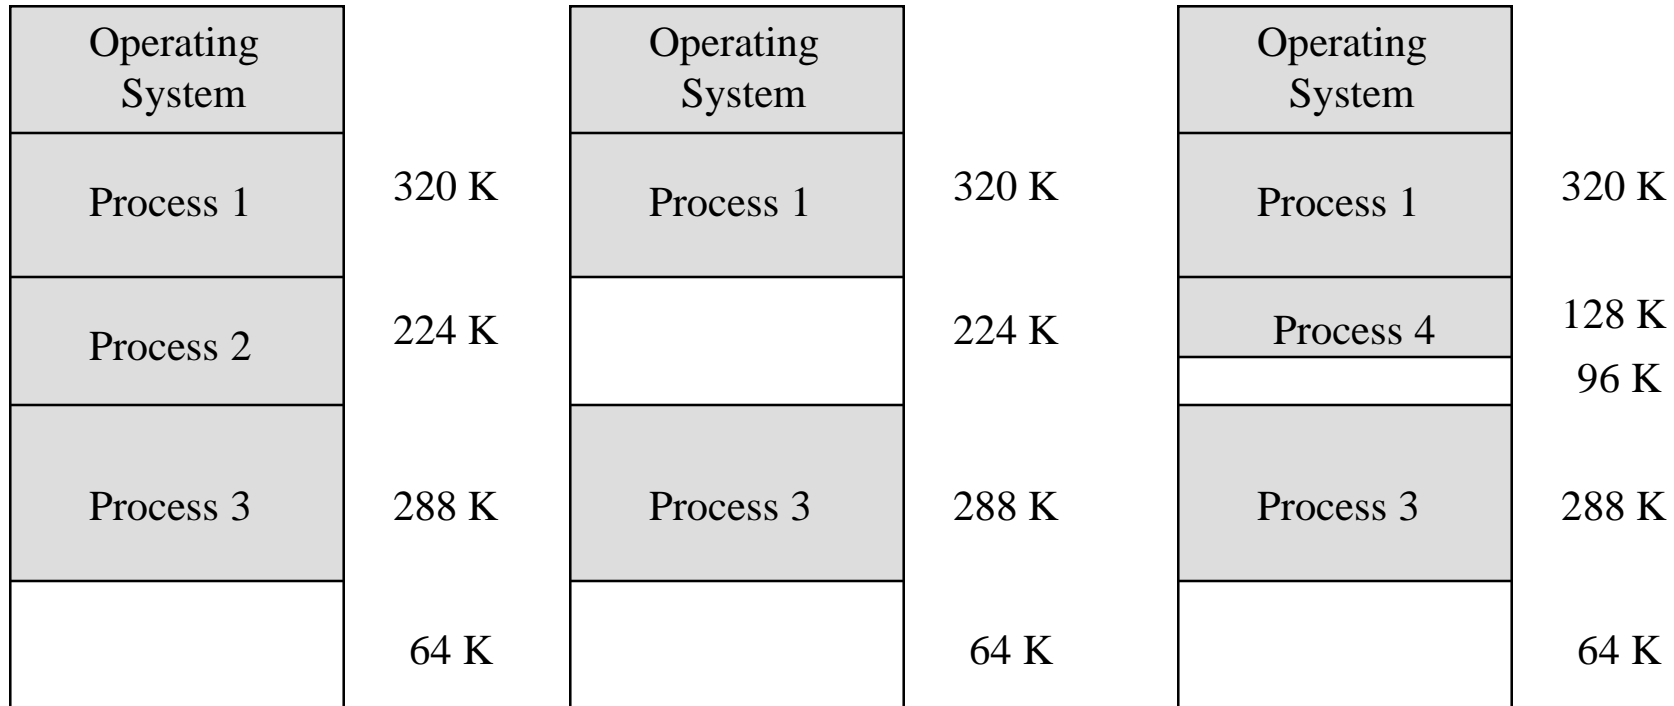
- ❑ Arrival and work requirements of 4 jobs
- ❑ CPU utilization for 1-4 jobs with 80% I/O wait
- ❑ Sequence of events as jobs arrive and finish
 - Numbers show amount of CPU time jobs get in each interval
 - More processes => better utilization, less time per process



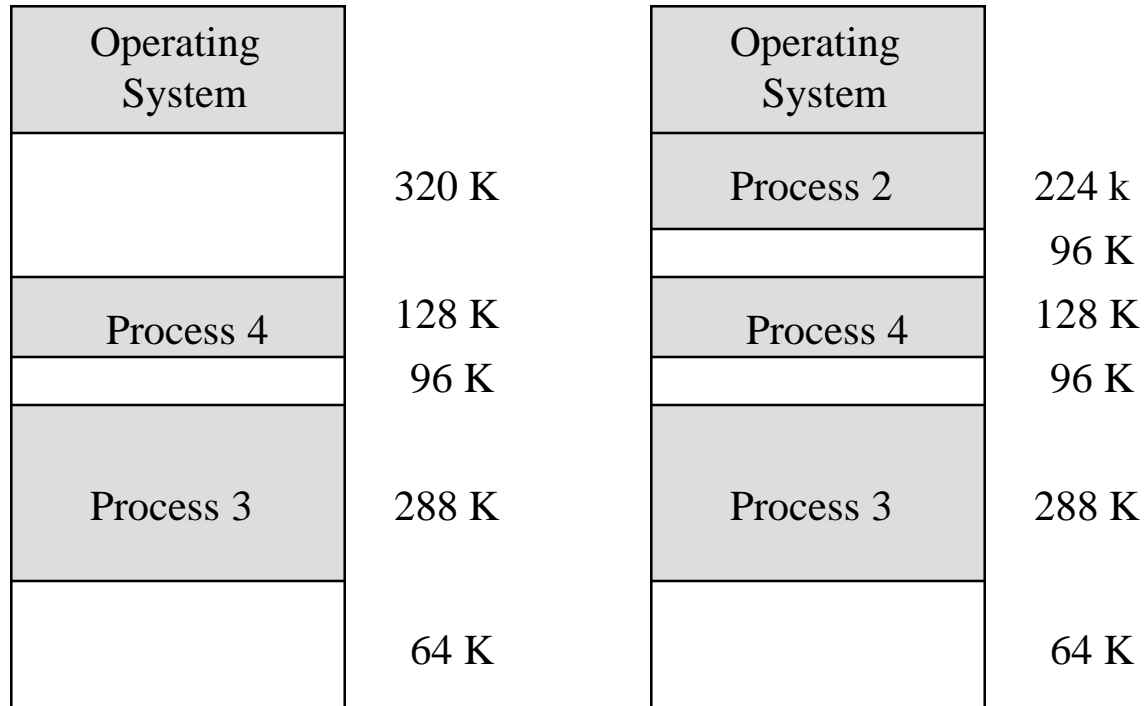
Dynamic Partitioning

- ❑ Partitions are of variable length and number
- ❑ Process is allocated exactly as much memory as required
- ❑ Eventually get holes in the memory. This is called **external fragmentation**
- ❑ **Compaction** is required to obtain a large block at the end of memory
 - Shift processes so they are contiguous and all free memory is in one block

Example Dynamic Partitioning



Example Dynamic Partitioning



Dynamic Partitioning Placement Algorithm

- ❑ Operating system must decide which free block to allocate to a process
- ❑ **Best-fit** algorithm
 - Choose the block that is closest in size to the request
 - This has the worst overall performance
 - The smallest block is found for a process
 - The smallest amount of fragmentation is left;
 - Memory compaction must be done more often

Dynamic Partitioning Placement Algorithm

❑ First-fit algorithm

- Starts scanning memory from the beginning and chooses the first available block that is large enough.

❑ Next-fit

- Starts scanning memory from the location of the last placement and chooses the next available block that is large enough

Memory and multiprogramming

- ❑ Memory needs two things for multiprogramming
 - Relocation
 - Protection
- ❑ The OS cannot be certain where a program will be loaded in memory
 - Variables and procedures can't use absolute locations in memory
 - Several ways to guarantee this
- ❑ The OS must keep processes' memory separate
 - Protect a process from other processes reading or modifying its own memory
 - Protect a process from modifying its own memory in undesirable ways (such as writing to program code)

Registers Used during Execution

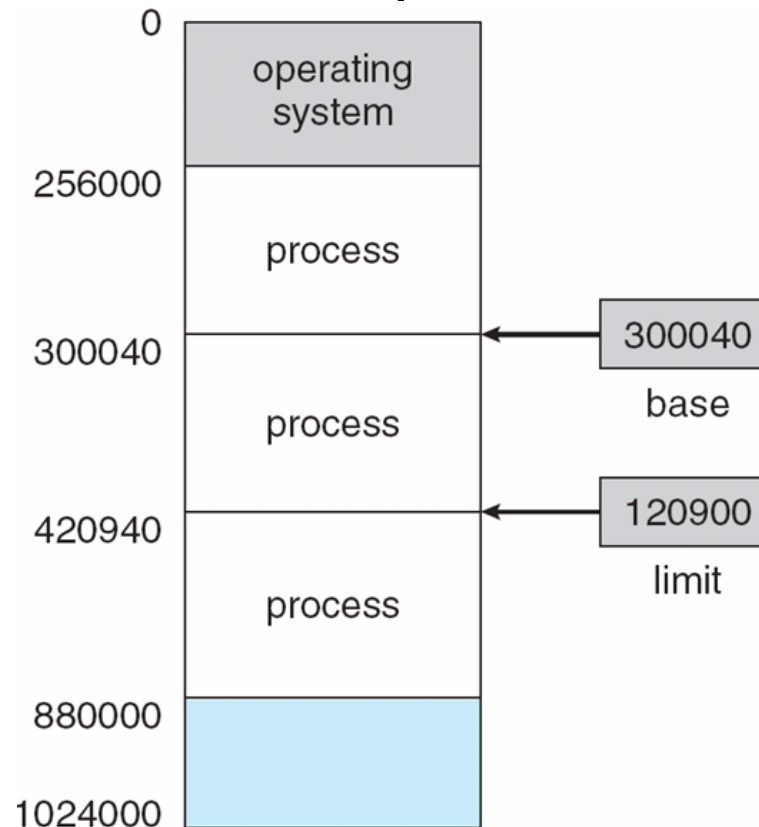
- ❑ Base register
 - Starting address for the process
- ❑ Bounds register
 - Ending location of the process
- ❑ These values are set when the process is loaded and when the process is swapped in

Registers Used during Execution

- ❑ The value of the base register is added to a relative address to produce an absolute address
- ❑ The resulting address is compared with the value in the bounds register
- ❑ If the address is not within bounds, an interrupt is generated to the operating system

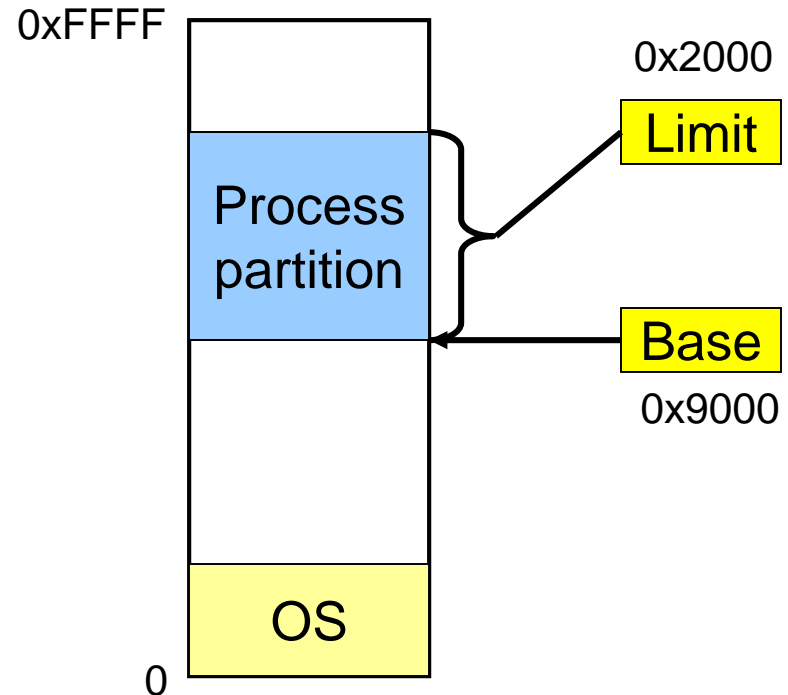
Base and Limit Registers

- A pair of **base** and **limit** registers define the logical address space



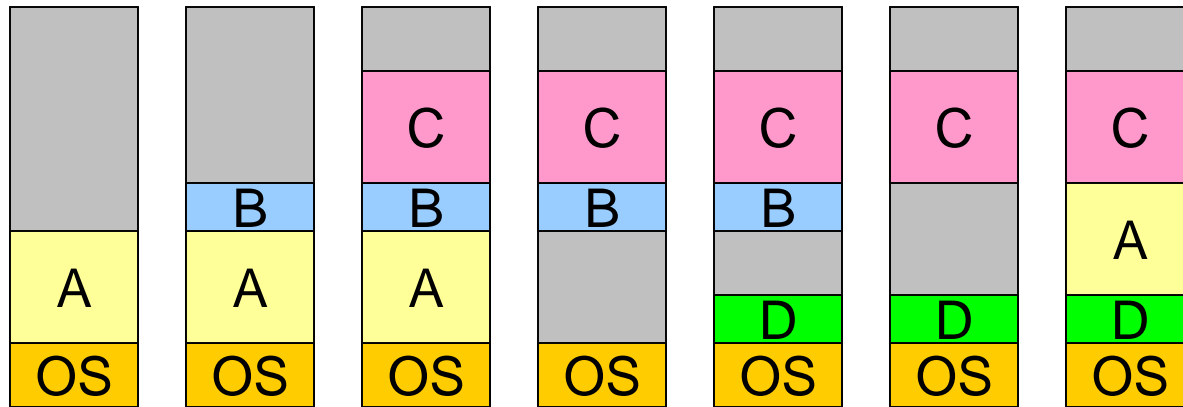
Base and limit registers

- ❑ Special CPU registers: base & limit
 - Access to the registers limited to system mode
 - Registers contain
 - Base: start of the process's memory partition
 - Limit: length of the process's memory partition
- ❑ Address generation
 - Physical address: location in actual memory
 - Logical address: location from the process's point of view
 - $\text{Physical address} = \text{base} + \text{logical address}$
 - Logical address larger than limit \Rightarrow error



Logical address: 0x1204
Physical address:
 $0x1204 + 0x9000 = 0xA204$

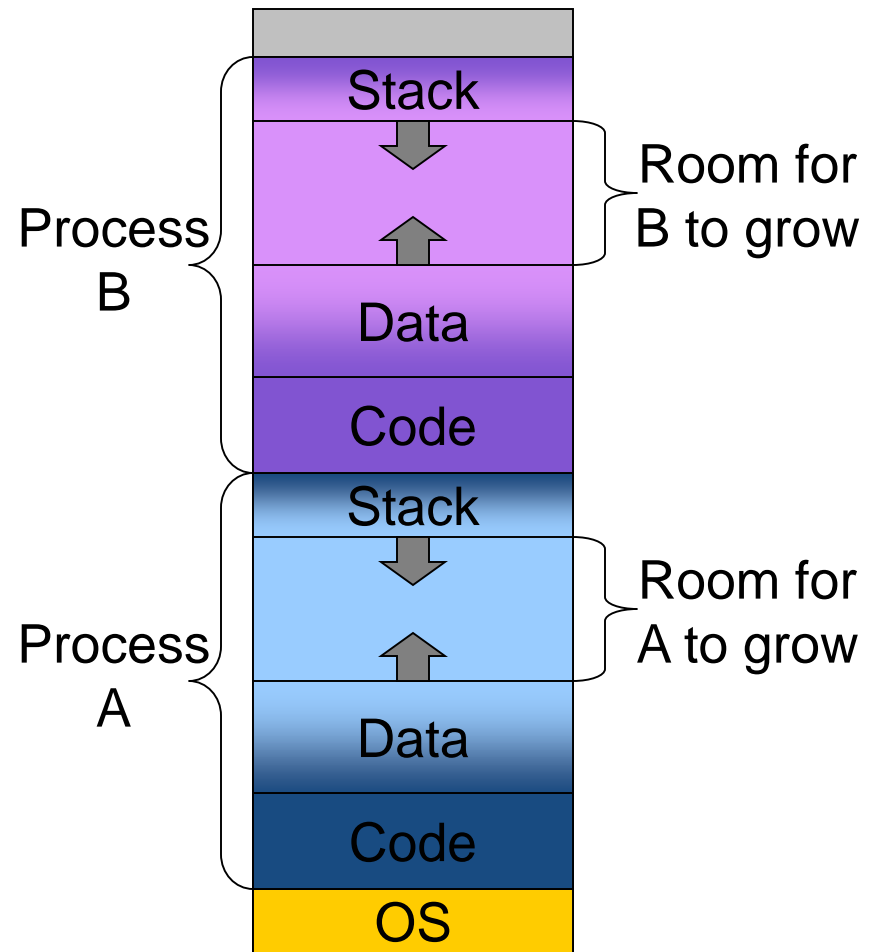
Swapping



- ❑ Memory allocation changes as
 - Processes come into memory
 - Processes leave memory
 - Swapped to disk
 - Complete execution
- ❑ Gray regions are unused memory


Swapping: leaving room to grow

- ❑ Need to allow for programs to grow
 - Allocate more memory for data
 - Larger stack
- ❑ Handled by allocating more space than is necessary at the start
 - Inefficient: wastes memory that's not currently in use
 - What if the process requests too much memory?



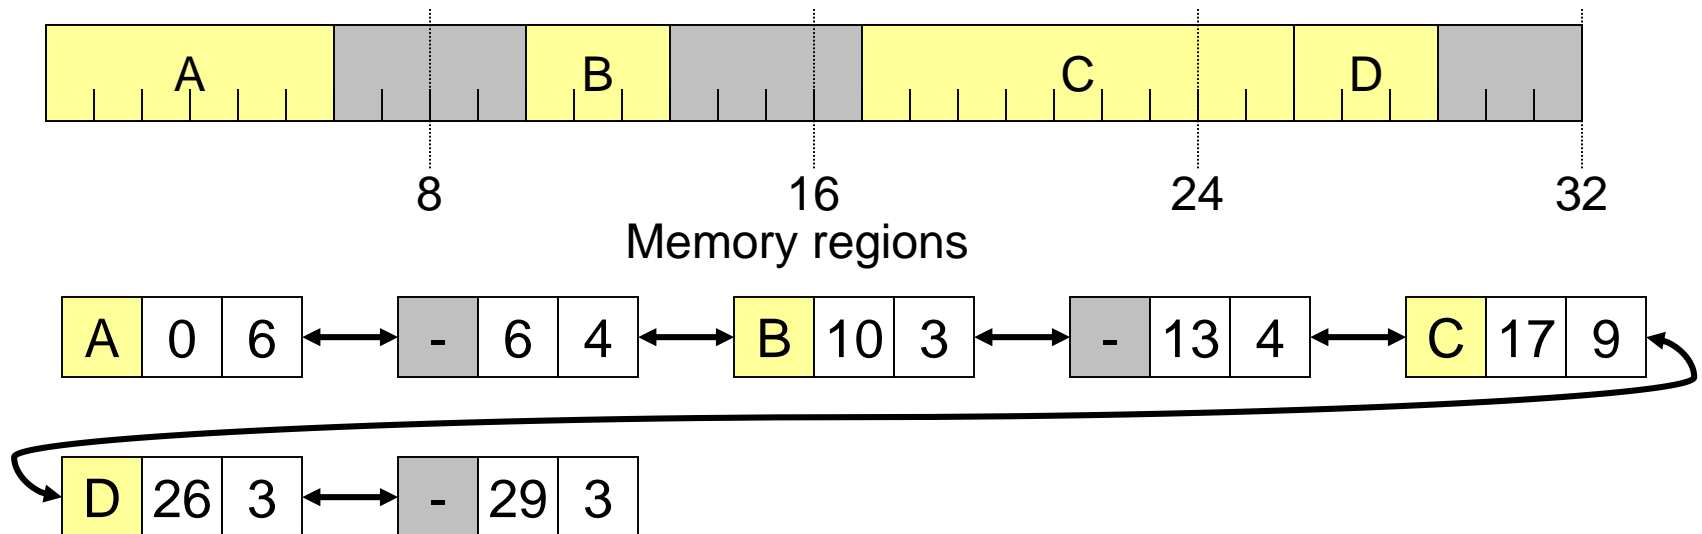
Tracking memory usage: bitmaps

- Keep track of free / allocated memory regions with a bitmap
 - One bit in map corresponds to a fixed-size region of memory
 - Bitmap is a constant size for a given amount of memory regardless of how much is allocated at a particular time
- Chunk size determines efficiency
 - At 1 bit per 4KB chunk, we need just 256 bits (32 bytes) per MB of memory
 - For smaller chunks, we need more memory for the bitmap
 - Can be difficult to find large contiguous free areas in bitmap



Tracking memory usage: linked lists

- Keep track of free / allocated memory regions with a linked list
 - Each entry in the list corresponds to a contiguous region of memory
 - Entry can indicate either allocated or free (and, optionally, owning process)
 - May have separate lists for free and allocated areas
- Efficient if chunks are large
 - Fixed-size representation for each region
 - More regions => more space needed for free lists




Allocating memory

- ❑ Search through region list to find a large enough space
- ❑ Suppose there are several choices: which one to use?
 - First fit: the first suitable hole on the list
 - Next fit: the first suitable after the previously allocated hole
 - Best fit: the smallest hole that is larger than the desired region (wastes least space?)
 - Worst fit: the largest available hole (leaves largest fragment)
- ❑ Option: maintain separate queues for different-size holes


Allocate 20 blocks first fit Allocate 13 blocks best fit

Allocate 12 blocks next fit Allocate 15 blocks worst fit



Freeing memory

- ❑ Allocation structures must be updated when memory is freed
- ❑ Easy with bitmaps: just set the appropriate bits in the bitmap
- ❑ Linked lists: modify adjacent elements as needed
 - Merge adjacent free regions into a single region
 - May involve merging two regions with the just-freed area



Limitations of swapping

❑ Problems with swapping

- Process must fit into physical memory (impossible to run larger processes)
- Memory becomes fragmented
 - External fragmentation: lots of small free areas
 - Compaction needed to reassemble larger free areas
- Processes are either in memory or on disk: half and half doesn't do any good

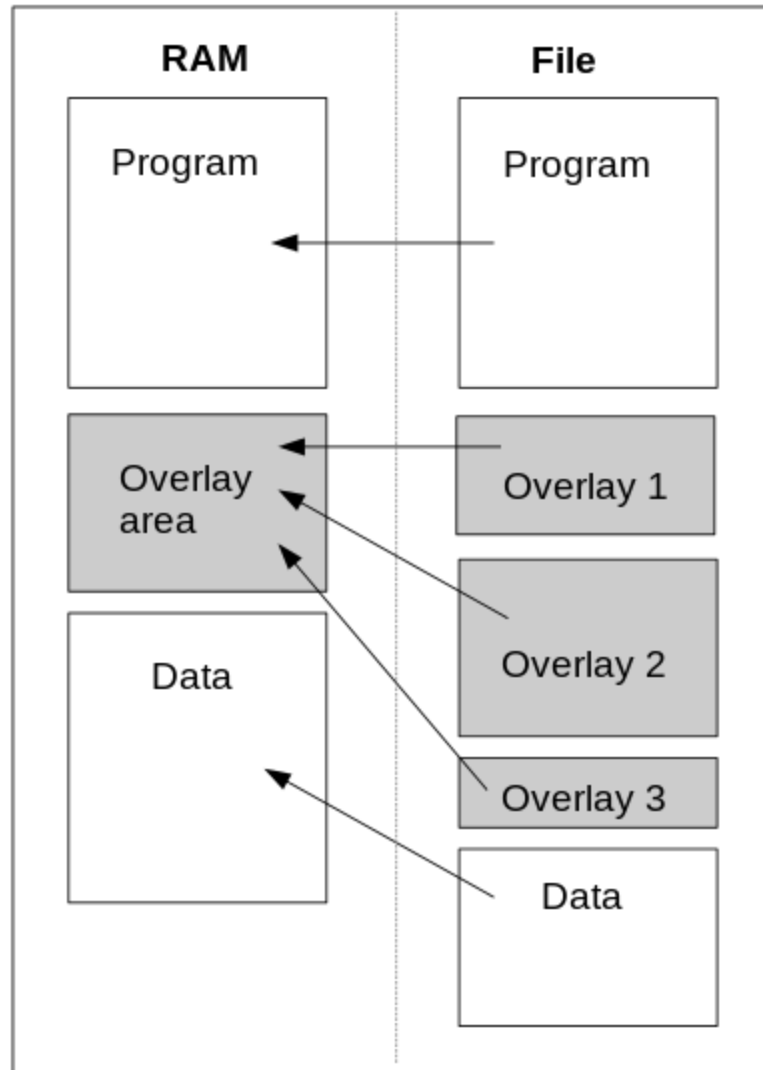
❑ Overlays solved the first problem

- Bring in pieces of the process over time (typically data)
- Still doesn't solve the problem of fragmentation or partially resident processes

Overlays

- ❑ Keep in memory only those instructions and data that are needed at any given time.
- ❑ Needed when process is larger than amount of memory allocated to it.
- ❑ Implemented by user, no special support from the operating system.
- ❑ Programming design of overlay structure is complex.

Overlays



Summary

- This section studied basic memory allocation techniques