

Introduction to Paging

Outline

- Introduction to Paging Concepts


Paging

- ❑ Partition memory into small equal-size chunks and divide each process into the same size chunks
- ❑ The chunks of a process are called **pages** and chunks of memory are called **frames**
- ❑ Operating system maintains a **page table** for each process
 - Contains the frame location for each page in the process
 - Memory address consist of a page number and offset within the page

Virtual memory

- ❑ Basic idea: allow the OS to hand out more memory than exists on the system.
- ❑ Keep recently used stuff in physical memory
- ❑ Move less recently used stuff to disk
- ❑ Keep all of this hidden from processes
 - Processes still see an address space from 0 - max address
 - Movement of information to and from disk handled by the OS without process help
- ❑ Virtual memory (VM) especially helpful in multiprogrammed system
 - CPU schedules process B while process A waits for its memory to be retrieved from disk


Virtual and physical addresses



- ❑ Program uses *virtual addresses*
 - Addresses local to the process
 - Hardware translates virtual address to *physical address*
- ❑ Translation done by the **Memory Management Unit**
 - Usually on the same chip as the CPU
 - Only physical addresses leave the CPU/MMU chip
- ❑ Physical memory indexed by physical addresses

Paging and page tables

- ❑ Virtual addresses mapped to physical addresses
 - Unit of mapping is called a *page*
 - All addresses in the same virtual page are in the same physical page
 - *Page table entry* (PTE) contains translation for a single page
- ❑ Table translates virtual page number to physical page number
 - Not all virtual memory has a physical page
 - Not every physical page need be used
- ❑ Example:
 - 64 KB virtual memory
 - 32 KB physical memory



Page Example

Frame
Number

0	0
1	1
2	2
3	3

Process A

0	7
1	8
2	9

Process B


0	A0
1	A1
2	A2
3	A3
4	
5	
6	
7	B0
8	B1
9	B2
10	
11	
12	
13	
14	

Paging

- ❑ A program requires N free frames
- ❑ There is a need to set up a page table to translate logical to physical addresses
- ❑ Internal fragmentation

Address Translation Scheme

- Address generated by CPU is divided into:
 - Page number (p) - used as an index into a page table which contains base address of each page in physical memory



- Page offset (d) - combined with base address to define the physical memory address that is sent to the memory unit
- The number of logical address bits is m

Address Translation Scheme

- ❑ If the number of logical address bits is m and the number of bits in the offset is n then
 - ❑ Logical address space is 2^m
 - ❑ Page size (number of entries) is 2^n
- ❑ Example: $m = 4, n = 2$
 - ❑ Number of address is 16
 - ❑ Number of entries in a page is 4

Mapping logical => physical address

- ❑ Split address from CPU into two pieces
 - Page number (p)
 - Page offset (d)
- ❑ Page number
 - Index into page table
 - Page table contains base address of page in physical memory
- ❑ Page offset
 - Added to base address to get actual physical memory address
- ❑ Page size = 2^d bytes


Example:

- 4 KB (=4096 byte) pages
- 32-bit logical addresses

$$2^d = 4096 \longrightarrow d = 12$$

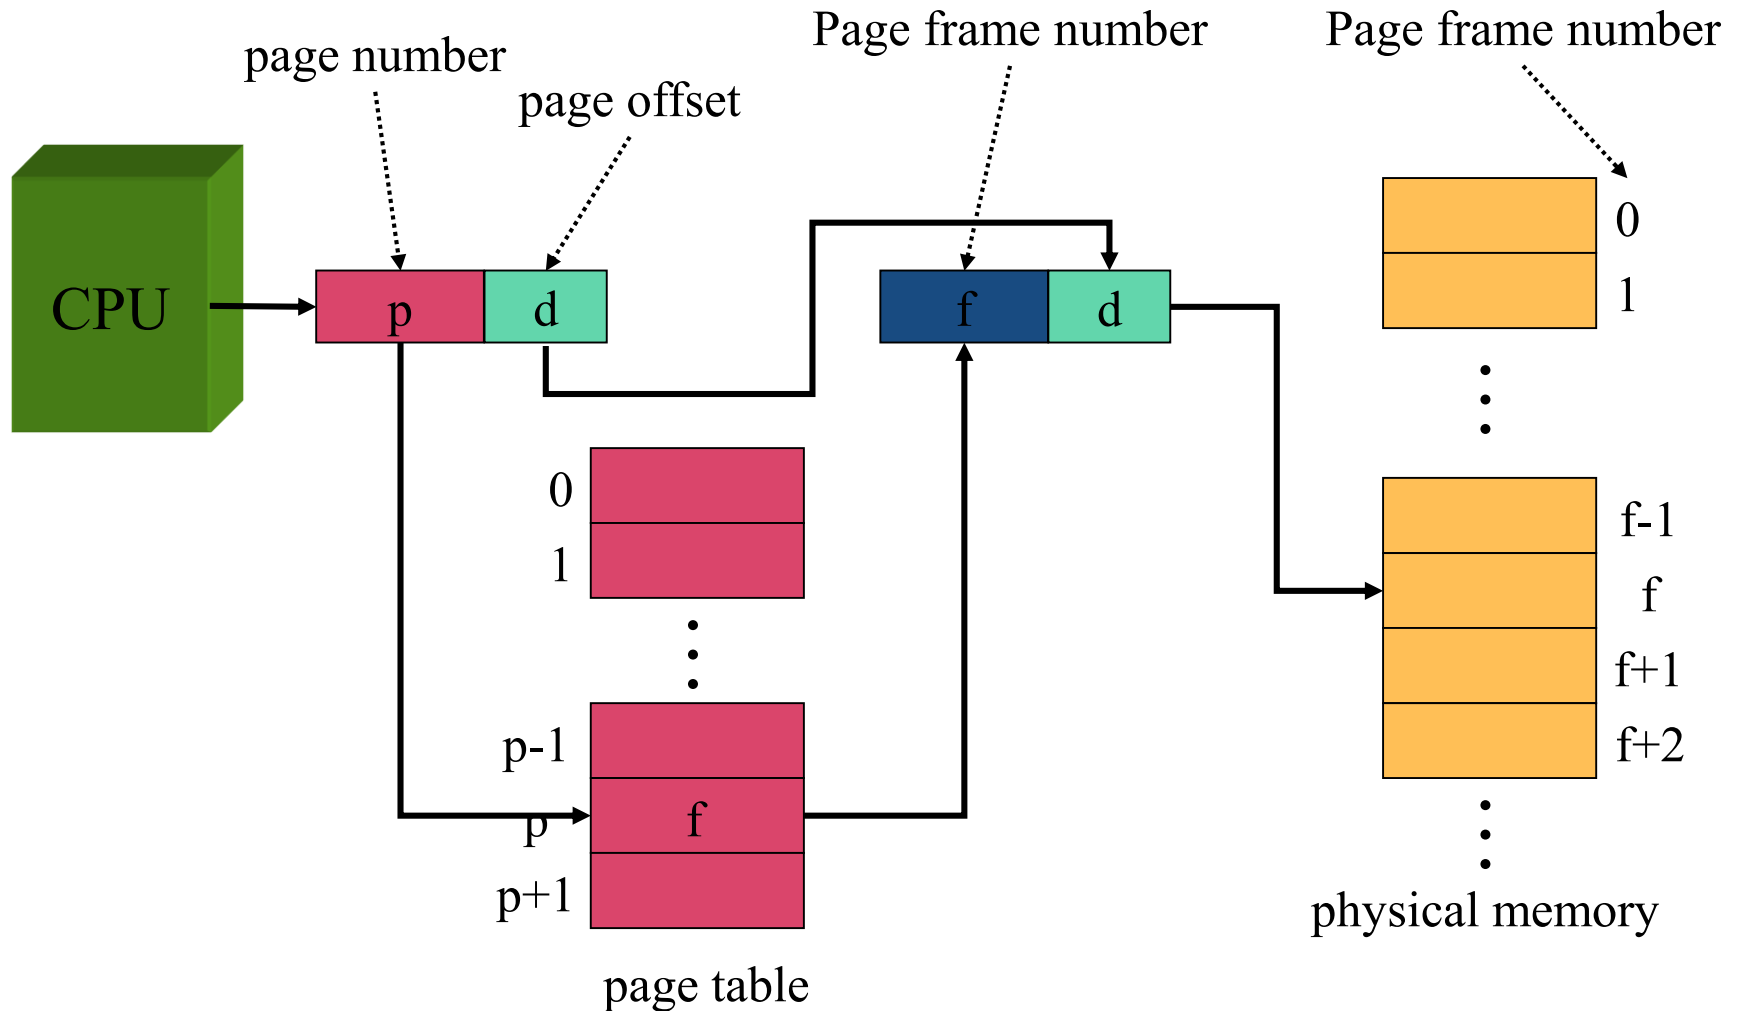


$$32 - 12 = 20 \text{ bits} \quad 12 \text{ bits}$$



32-bit logical address

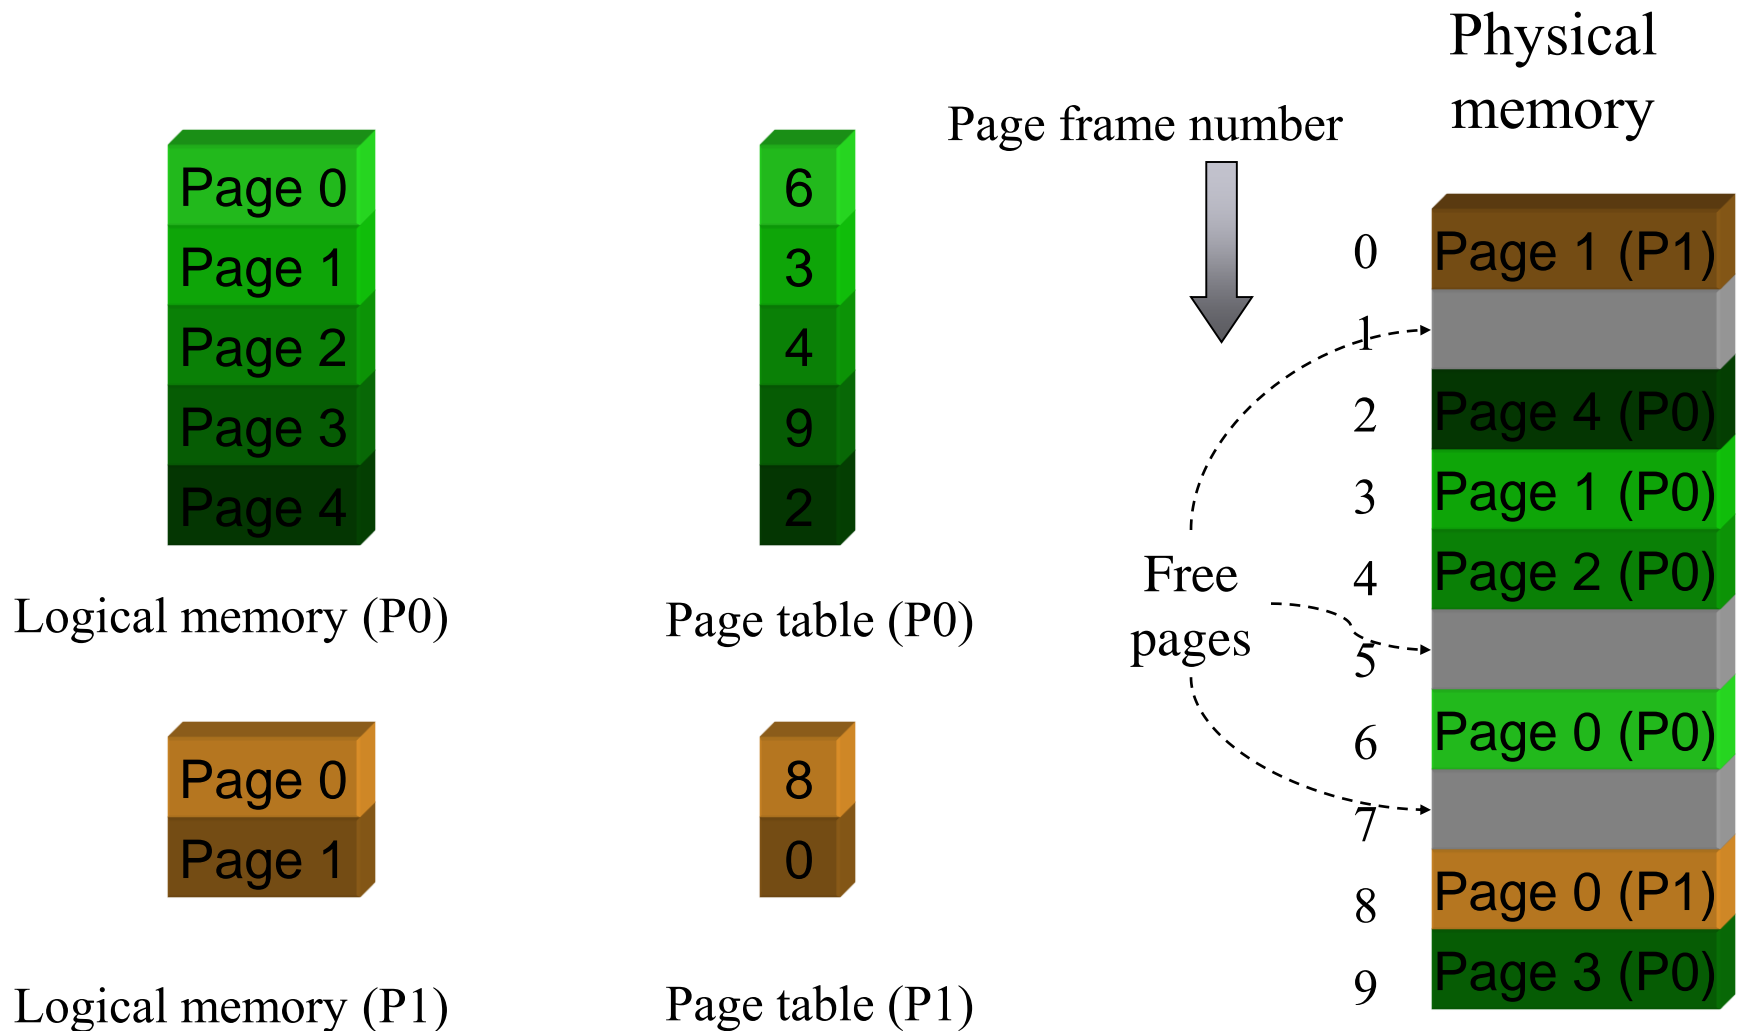
Address translation architecture



Paging Hardware

- ❑ The CPU issues a logical address (remember that all addresses are binary)
- ❑ The hardware extracts the page number, p , and the page offset d
- ❑ The page number, p , is used to index the page table
 - The entry in the page table consists of the frame number, f
- ❑ The actual address is the concatenation of the bits that make up f and d .

Memory & paging structures



Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

m=4

n=2

Page size of 4 bytes

Physical memory of 32 bytes (8 pages)

Binary representations:

0 00

1 01

2 10

3 11

Paging Example

- First a little reminder. All addresses are in bits. The example on the previous page gives the base 10 equivalent of binary numbers

0	0000	8	1000
1	0001	9	1001
2	0010	10	1010
3	0011	11	1011
4	0100	12	1100
5	0101	13	1101
6	0110	14	1110
7	0111	15	1111

Example

- ❑ Logical address in binary form is 0000
 - From the bits we extract the page number as 00 and the offset as 00
 - Page number 00 corresponds to the first entry of the page table.
 - There we find that the corresponding frame is 5 (or 0101)
 - The physical address produced is the concatenation of the bits of the frame number and the offset
 - 010100
 - In base 10 we find that the physical address is 20 ($5 \times 4 + 0$)

Example

- ❑ Logical address in binary form is 0100
 - From the bits we extract the page number as 01 and the offset as 00
 - Page number 01 corresponds to the second entry of the page table.
 - There we find that the corresponding frame is 6 (or 0110)
 - The physical address produced is the concatenation of the bits of the frame number and the offset
 - 011000
 - In base 10 we find that the physical address is $24(6*4+0)$
- ❑ What do you think logical address 13 maps to?

Observations

- Let's look at logical addresses:
 - 0000 (0), 0001 (1), 0010(2) and 0011(3)
 - The first two bits (page number) are the same
 - This means that these logical addresses will be in the same frame.
 - The position in the frame is determined by the offset
- For a process its pages can be in any order
 - For our example page 1 is in a frame that appears "later" than the frame associated with page 2

Observations

- ❑ The logical address space and the physical address space DO NOT have to be the same size
- ❑ The logical address space can be larger than the physical address space (more on this later)

Paging

- ❑ No external fragmentation
 - Any free frame can be allocated to a process that needs it
- ❑ There may be some internal fragmentation
 - The memory requirements may not coincide with page boundaries
 - The last frame allocated may not be completely full

Paging

- ❑ Internal fragmentation is on average one half page per process
- ❑ Larger page sizes means more wasted space
- ❑ Smaller page sizes means larger tables

Summary

- This section introduced the concept of paging in memory