

Page Table Implementation

Outline

- ❑ Structure of the Page Table
- ❑ Implementation Issues
- ❑ Segmentation



Frames

- ❑ Each process page needs one frame.
- ❑ The OS needs to be aware of the available frames.
- ❑ This is through the **frame table**
 - Each frame has an entry in the table
 - Each entry determines if the frame is allocated or not
 - If allocated indicates which page

Free Frames

free-frame list



14
13
18
20
15



(a)

Before allocation

free-frame list
15



(b)

0	14
1	13
2	18
3	20

new-process page table

After allocation

Implementation of Page Table

- ❑ Page table is kept in main memory
- ❑ **Page-table base register** (PTBR) points to the page table
- ❑ **Page-table length register** (PTLR) indicates size of the page table

Implementation of Page Table

- ❑ **Problem:** Every data/instruction access requires two memory accesses:
 - Access for the page table
 - Access for the data/instruction.
- ❑ **Solution:** Use a special fast-lookup **hardware cache** called associative memory or Translation Look-aside Buffers (TLBs)
- ❑ TLBs store address-space identifiers (ASIDs) in each TLB entry - uniquely identifies each process to provide address-space protection for that process


TLB

- ❑ Associative memory - parallel search
- ❑ Address translation (p, d)
 - If p is in associative register, get frame number out
 - Otherwise get frame number from page table in memory

Page #	Frame #

- ❑ Address-space identifiers (ASIDs) are associated with each entry
 - uniquely identifies each process to provide address-space protection for that process

Paging Hardware with TLB



Page Table with TLB (cont.)

- ❑ **TLB**: Cache for VM address to PM address lookup.
 - Like a dedicated cache for lookup.
 - TLB entries contain both virtual address and its corresponding physical address.
- ❑ **Page Table**: Map VM address to PM address as well, but it is slow.
- ❑ **hardware cache** called associative memory or Translation Look-aside Buffers (TLBs)
- ❑ TLBs store address-space identifiers (ASIDs) in each TLB entry - uniquely identifies each process to provide address-space protection for that process

Effective Access Time

- ❑ Associative Lookup = ε time unit
- ❑ Assume memory cycle time is 1 microsecond
- ❑ Hit ratio: percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- ❑ Hit ratio = α
- ❑ **Effective Access Time (EAT)**
- ❑ Example 1: 80% hits, 100 nanosecs to access memory and 20 nanosecs to search TLB.
Determine EAT?
- ❑ Example 2: 98% hits, determine EAT?

Effective Access Time (cont.)

- ❑ TLB hit time = TLB search time + memory access time
- ❑ TLB miss time = TLB search time + memory access time + memory access time.
- ❑ $EAT = TLB \text{ miss time} * (1 - \text{hit ratio}) + TLB \text{ hit time} * \text{hit ratio}$
- ❑ Example 1: $EAT = 0.80 * 120 + 0.20 * 220 = 140ns$
- ❑ Example 2: $EAT = 0.98 * 120 + 0.02 * 220 = 122ns$

Shared Pages


❑ Shared code

- One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).
- Shared code must appear in same location in the logical address space of all processes

❑ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example



Structure of the Page Table

- ❑ Typically, systems have large logical address spaces (2^{32} to 2^{64})
- ❑ Page table could be excessively large
- ❑ Example:
 - 32-bit logical address space
 - The number of possible addresses in the logical address space is 2^{32}
 - Page size is 4 KB (2^{12})
 - Page table may consist of up to 1 million entries ($2^{32} / 2^{12}$)
 - Assume each entry consists of 4 bytes
 - A process may need up to 4 MB for the page table


Structure of the Page Table

- ❑ The mapping should be fast
- ❑ Note that each memory reference requires access to the page table
- ❑ A typical instruction has an instruction word, and often a memory operand
- ❑ There is the potential for multiple page table references per instruction
- ❑ Why not have a single page table consisting of an array of fast hardware registers
 - Expensive in terms of registers used

Structure of Page Table

- ❑ Several approaches
 - Hierarchical Paging (multiple levels)
 - Hashed Page Tables
 - Inverted Page Tables

Two-Level Page-Table Scheme



Two-Level Paging Example

- ❑ A logical address (on 32-bit machine with 1K page size) is divided into:
 - A page number consisting of 22 bits
 - A page offset consisting of 10 bits
- ❑ Since the page table is paged, the page number is further divided into:
 - A 12-bit page number
 - A 10-bit page offset

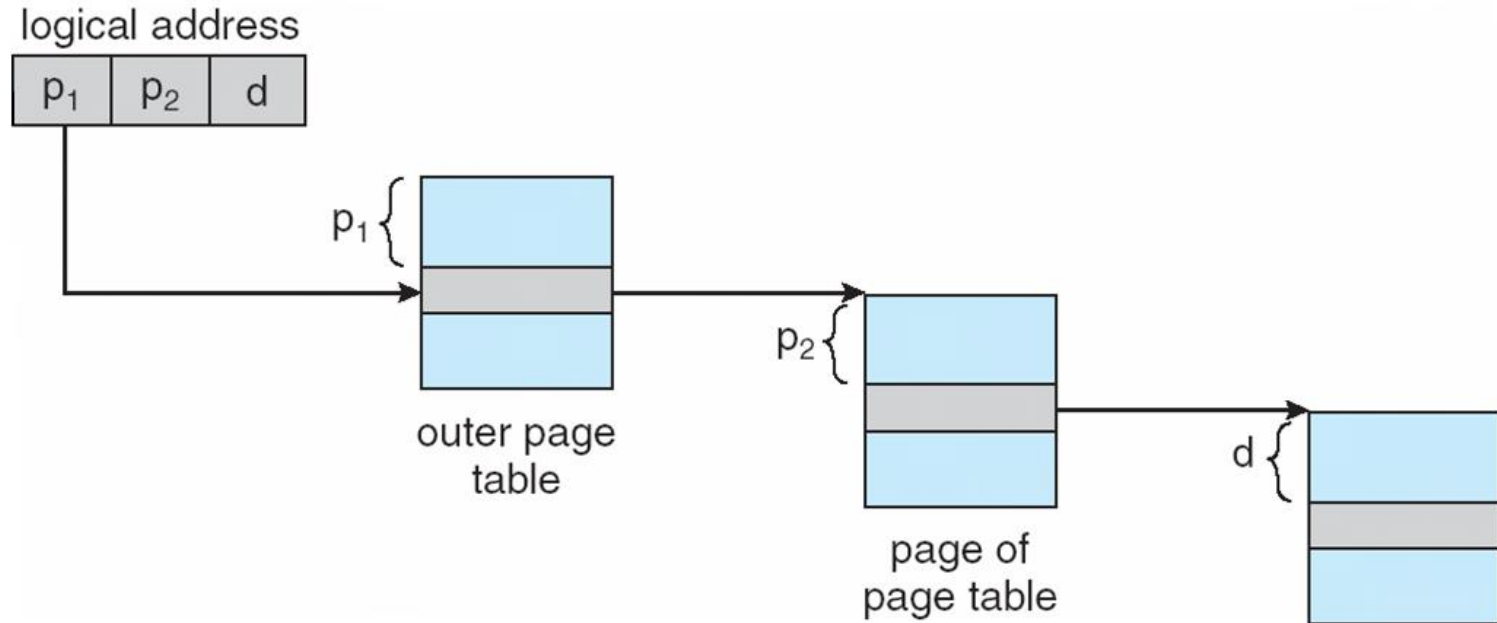
Two-Level Paging Example(cont.)

- Thus, a logical address is as follows:

page number		page offset
p_i	p_2	d
12	10	10



- Here p_i is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

Address-Translation Scheme



- p_1 is used to index the outer page table
- p_2 is used to index the page table

Three-level Paging Scheme




Hierarchical Paging

- ❑ What if you have a 64-bit architecture?
- ❑ Do you think hierarchical paging is a good idea?
 - How many levels of paging are needed?
 - What is the relationship between paging and memory accesses?

Hashed Page Tables

- ❑ Common in address spaces > 32 bits
- ❑ The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- ❑ Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted

Hashed Page Table



Inverted Page Table

- ❑ One entry for each real page of memory
- ❑ Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- ❑ Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- ❑ To address this problem, use a hash table
- ❑ Used in the 64-bit UltraSparc and PowerPC

Inverted Page Table Architecture

Logical address (p1)

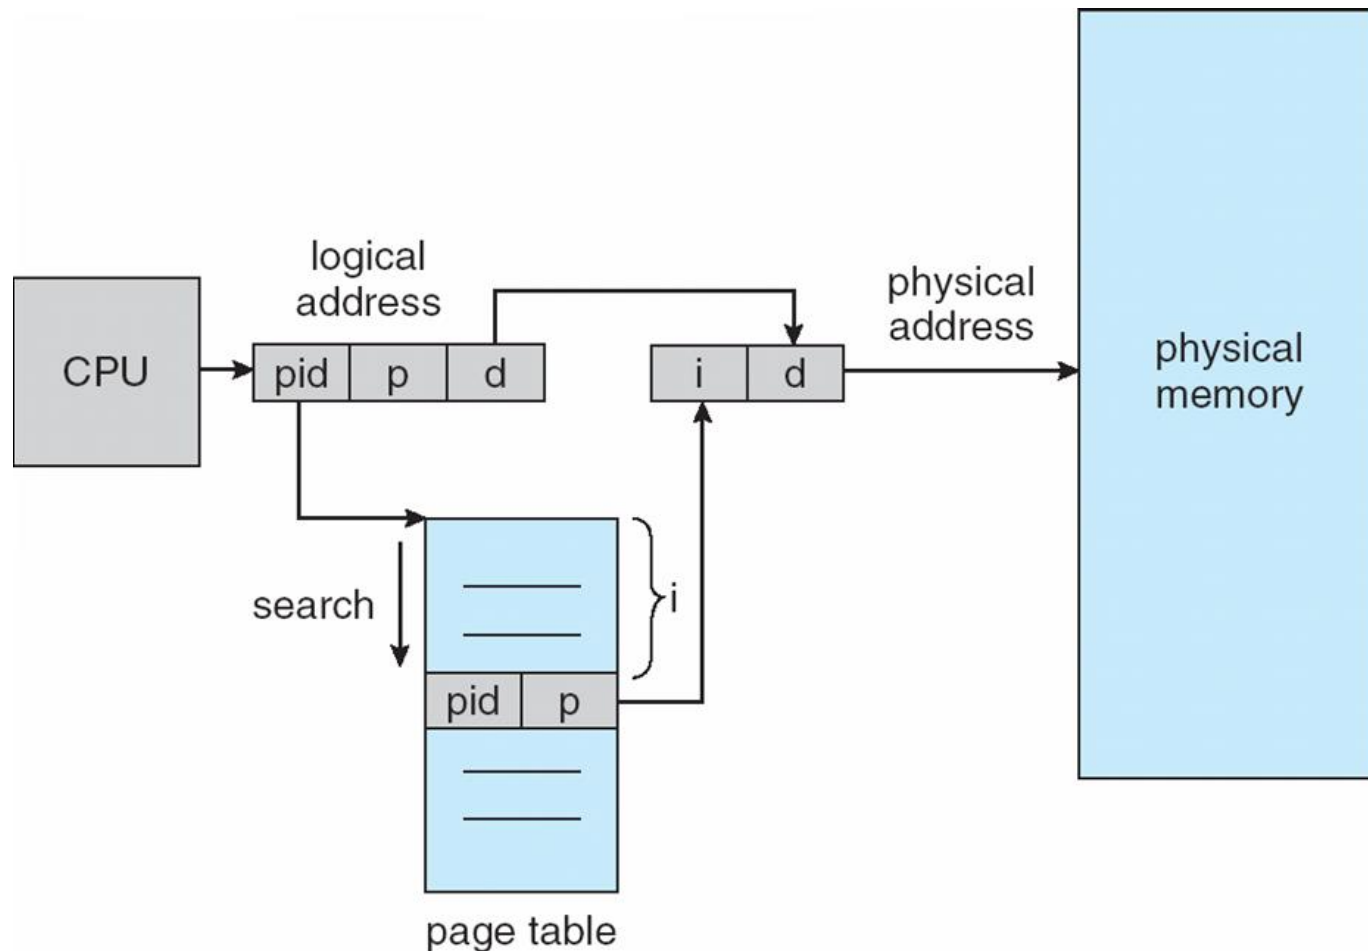
Page no | offset

Inverted Page Table

0		
1	Page0	p3
2	page0	p2
3	Page0	p1
4		
5	page3	p1
6	page4	p3

Page 0, process 1

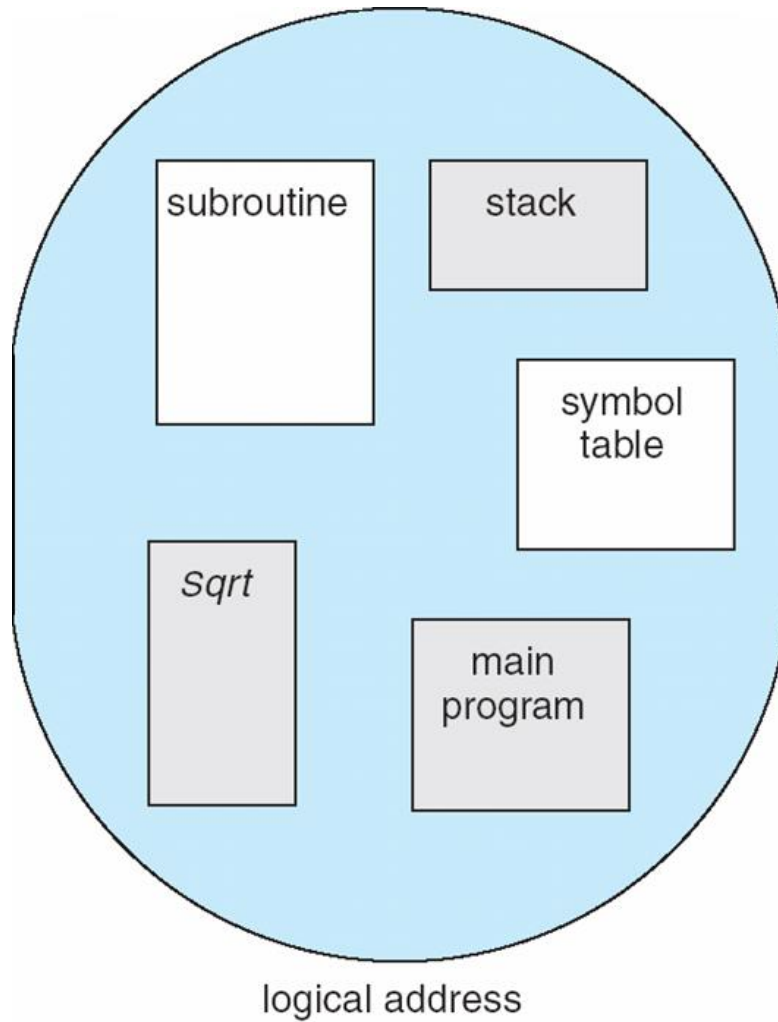
Inverted Page Table Architecture (cont.)



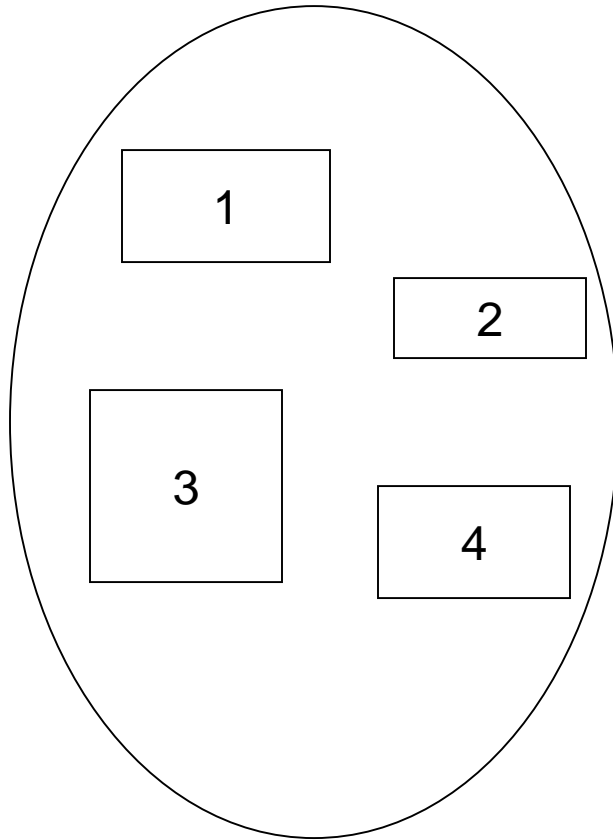
Segmentation

- ❑ Memory-management scheme that supports user view of memory
- ❑ A program is a collection of segments
 - A segment is a logical unit such as:
 - main program
 - procedure
 - function
 - method
 - object
 - local variables, global variables
 - common block
 - stack
 - symbol table
 - arrays


User's View of a Program



Logical View of Segmentation



user space



physical memory space

Segmentation Architecture

- ❑ Logical address consists of a two tuple:
 <segment-number, offset>,
- ❑ **Segment table** - maps two-dimensional physical addresses; each table entry has:
 - **base** - contains the starting physical address where the segments reside in memory
 - **limit** - specifies the length of the segment
- ❑ **Segment-table base register (STBR)** points to the segment table's location in memory
- ❑ **Segment-table length register (STLR)** indicates number of segments used by a program;
 segment number **s** is legal if **s** < **STLR**

Segmentation Architecture (Cont.)

❑ Protection

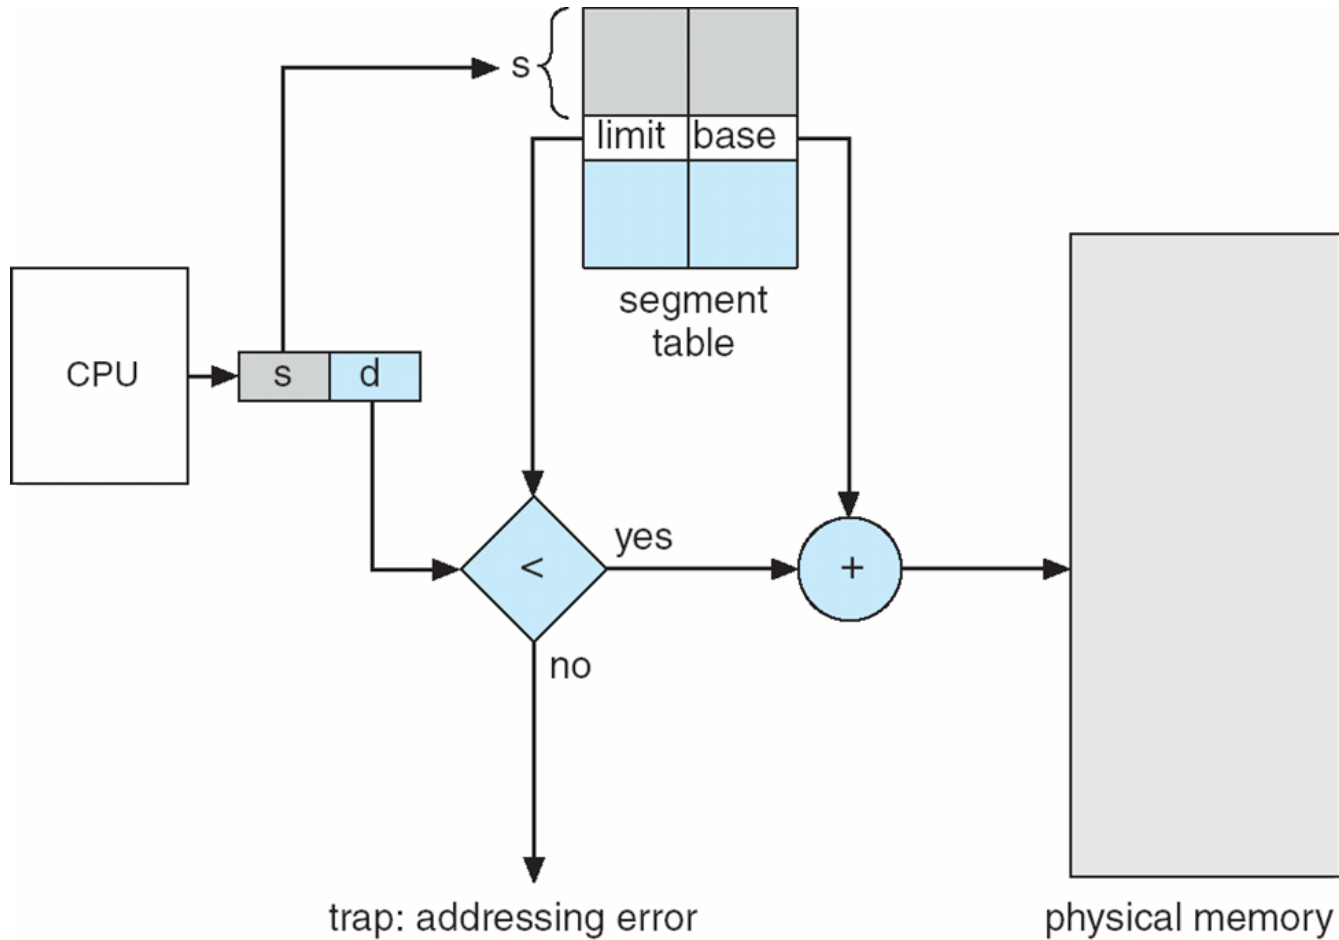
- With each entry in segment table associate:
 - validation bit = 0 \Rightarrow illegal segment
 - read/write/execute privileges

❑ Protection bits associated with segments; code sharing occurs at segment level


❑ Since segments vary in length, memory allocation is a dynamic storage-allocation problem

❑ A segmentation example is shown in the following diagram

Segmentation Hardware




Example of Segmentation



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table



Summary

- This section studied how page tables are implemented. It also introduced the concept of segmentation.