

Virtual Memory

Program

```
# include <stdio.h>

....

main()
{
    /*Lots and lots of initialization code*/

    for (i = 0; i < N; i++)
    {
        ...
    }

    ...
}
```

Program

- ❑ Typically once you initialize you don't go back to that part of the code.
- ❑ Do you need to keep that part of the code in main memory once it has been used?

Virtual Memory: Main Idea

- ❑ Processes use a virtual (logical) address space.
- ❑ Every process has its own address space
- ❑ The virtual address space can be larger than physical memory.
- ❑ Only part of the virtual address space is **mapped** to physical memory at any time.
- ❑ Parts of processes' memory content is on disk.
- ❑ Hardware & OS collaborate to move memory contents to and from disk.


Demand Paging

- ❑ Bring a page into memory only when it is needed
 - Why? Less I/O needed
 - If a process of 10 pages actually uses only half of them, then demand paging saves the I/O necessary to load the 5 pages not used.
 - Less memory needed
 - Faster response
 - More multiprogramming is possible

Demand Paging

- ❑ We need hardware support to distinguish between pages that are in memory and the pages that are on disk
- ❑ A valid-invalid bit is part of each page entry
 - When the bit is set to "valid" the associated page is in memory
 - If the bit is set to "invalid" the page is on the disk

Demand Paging

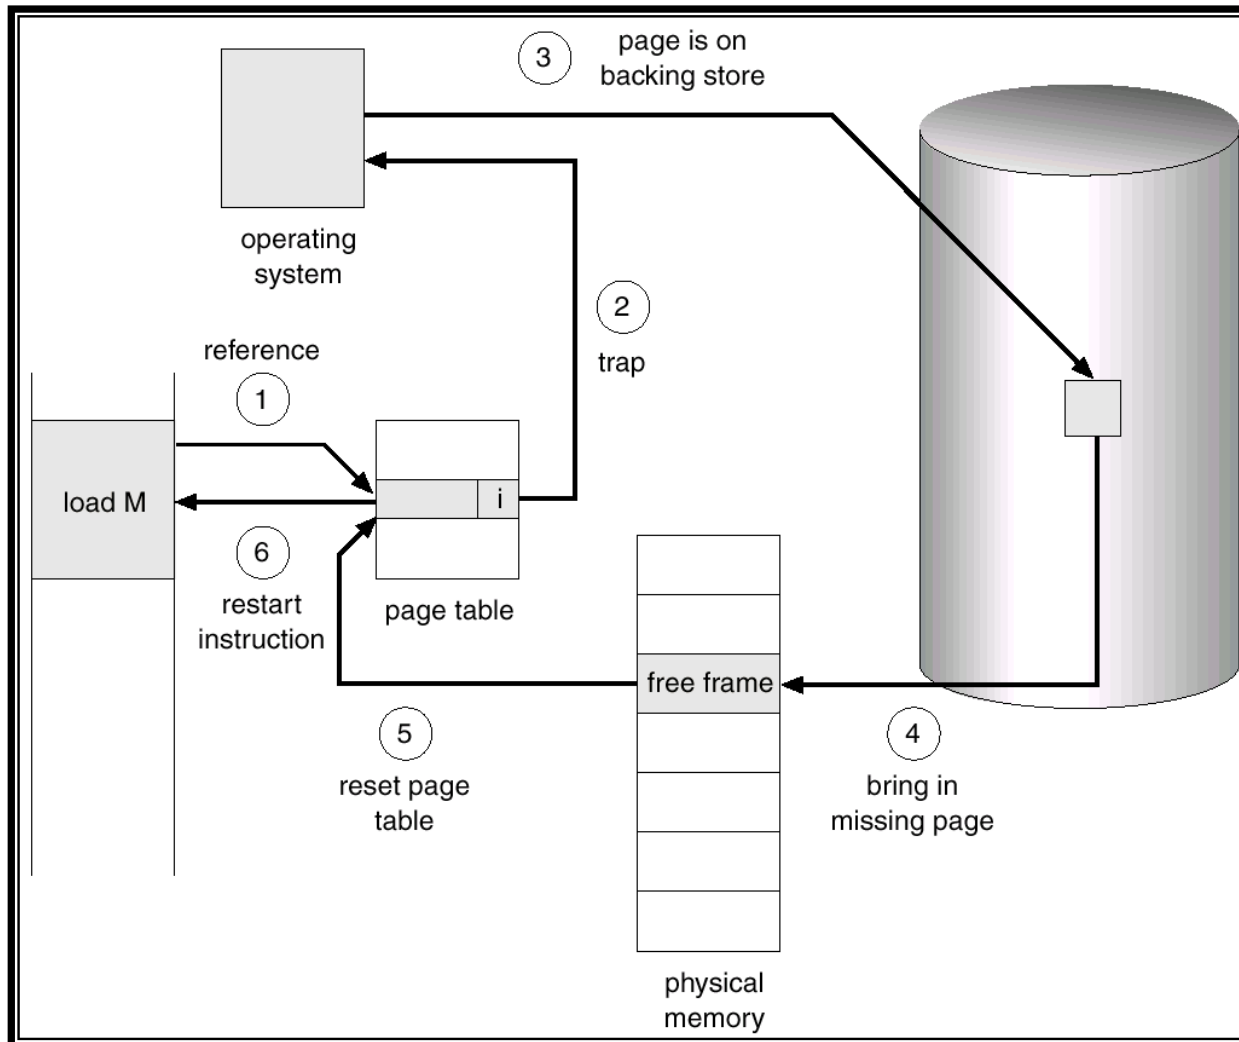


- The valid-invalid bit for 1 is set to “i” since the page is not in the physical memory
- The valid-invalid bit for 0 is “v” since the page is in memory

Page Fault

- ❑ What happens if a process tries to access a page that was not brought into memory?
 - Example: User opens a Powerpoint file.
- ❑ Access to a page marked invalid causes a **page fault**
- ❑ The paging hardware, in translating the address through the page table, will notice that the invalid bit is set causing a trap to the operating system.

Steps in Handling a Page Fault



Steps in Handling a Page Fault

- ❑ Service the page-fault interrupt
 - Trap to the OS
 - Prepare for a context switch
 - Save the user registers and process state
 - Read in the page
 - Issue a read from the disk to a free frame
 - Wait in the queue for this device until the read request is serviced
 - Wait for the device seek and/or latency time
 - Begin the transfer of the page to a free frame

Steps in Handling a Page Fault

- ❑ While waiting, allocate the CPU to some other user
- ❑ Receive an interrupt from the disk I/O subsystem

Steps in Handling a Page Fault

- ❑ Deal with Interrupt from the disk I/O system
 - Save the registers and process state for the other user
 - Correct the page table
 - Process with a page fault is put into the ready queue
- ❑ Wait for the CPU to be allocated to the process again
- ❑ Restore the user registers, process state and new page table, and then resume the interrupted instruction

Challenge: Performance

- ❑ Page Fault Rate $0 \leq p \leq 1.0$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- ❑ Let p be the probability of page fault:
 - Avg. time = $(1-p) * \text{memory time} + p * \text{page fault time}$
 - Assuming: memory time = 200ns, page fault time = 8 millisecond, $p = 0.1\%$
 - Avg time = $99.9\% * 200 + 0.1\% * 8000000 = 8200$
 - Access time is directly proportional to the probability of a page fault
 - If one access out of 1000 causes a page fault the effective access time is 8.2 microseconds
- ❑ Need to keep the page fault rate small!

Page Replacement

- ❑ So why allow demand paging?
- ❑ If a process of 10 pages actually uses only half of them, the demand paging saves the I/O necessary to load the 5 pages that are never used
- ❑ This allows us to increase the level of multiprogramming

Page Replacement

- ❑ Let's assume that our physical memory consists of 40 frames
- ❑ We have 8 processes with 10 pages. That is 80 pages.
 - Obviously 80 pages is more than 40 frames
 - But if a process is only using half of its pages is this really a problem?
- ❑ But there is a reason why there are 10 pages
 - The process may need them
- ❑ The frames have been **over-allocated** i.e., overbooked

Page Replacement

- ❑ What do we do when a process needs a frame and there isn't one free?
- ❑ Essentially we choose a frame and free it of the page that is currently residing on it

Page Replacement

- ❑ A page replacement algorithm describes which frame becomes a victim.
- ❑ Designing an appropriate algorithm is important since disk I/O is expensive
- ❑ Slight improvements in algorithms yield large gains in system performance

Page Replacement

- ❑ We will discuss several algorithms
- ❑ The examples assume:
 - 3 frames
 - Reference string:
7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1
 - Each of the numbers refers to a page number

Optimal Page Replacement Algorithm

- ❑ Replace page needed at the farthest point in future i.e. replace the page that will not be used for the longest period of time
- ❑ This should have the lowest page fault rate

Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2			2			2				7		
	0	0	0		0		4			0			0				0		
		1	1		3		3			3			1				1		

page frames

Optimal Page Replacement

- ❑ Optimal is easy to describe but impossible to implement
- ❑ At the time of the page fault, the OS has no way of knowing when each of the pages will be referenced next

FIFO Page Replacement Algorithm

- ❑ Maintain a linked list of all pages
 - Each page is associated with the time when that page was brought into memory
- ❑ Page chosen to be replaced is the oldest page
- ❑ Implementation: FIFO queue
 - A variable **head** points to the oldest page
 - A variable **tail** points to the newest page brought in

FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																	
	0	0	0																	
		1	1																	

2	2	4	4	4	0															
3	3	3	2	2	2															
1	0	0	0	3	3															

0	0																			
1	1																			
3	2																			

7	7	7																		
1	0	0																		
2	2	1																		

page frames

FIFO Page Replacement

❑ Advantages

- Easy to understand and program

❑ Disadvantage

- Performance is not always good
- The page replaced may be an initialization module that was used a long time ago and is no longer needed, **but on the other hand ...**
 - The page may contain a heavily used variable that was initialized early and is in constant use

LRU Page Replacement

- ❑ FIFO replacement algorithm uses time when a page was brought into memory
- ❑ The optimal replacement algorithm uses the time when a page is to be used.
- ❑ Can we use the recent past as an approximation of the near future?
 - This means replace the page *that has not been used* for the longest period of time
 - This approach is **the Least-Recently-Used (LRU)** algorithm.

LRU Replacement Algorithm

- ❑ LRU replacement associates with each page the time that of that page's last use
- ❑ When a page must be replaced, LRU chooses the page that has not been used for the longest period of time.

LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

LRU Page Replacement

- ❑ LRU is often used and is considered to be good
- ❑ Challenge: Implementing LRU
- ❑ Implementation using Counters:
 - Associate with each page-table entry a time-of-use field
 - Add to the CPU a logical clock or counter
 - Clock is incremented for every memory reference
 - Each time a page is referenced the time-of-use field is updated with the logical clock
 - A search of the page table is needed to find the least recently used page

LRU Page Replacement

- ❑ Implementation using Stack:
 - Keep a stack of page numbers
 - When a page is referenced, it is removed from the stack and put on the top
 - The most recently used page is always at the top of the stack and the least recently used is at the bottom
 - Should use a doubly-linked list since entries can be removed from the middle of the stack

LRU Approximation Algorithms

❑ LRU needs special hardware and still slow

❑ Reference bit

- With each page associate a bit, initially = 0
- When page is referenced, bit set to 1
- Replace any with reference bit = 0 (if one exists)
 - We do not know the order, however

LRU Page Replacement

❑ Implementation using Reference Bits:

- Each page entry has a set of reference bits e.g., 8 bits
- The bits represent the history of page use for the last 8 time periods
 - 00000000: implies that the page has not be used in the last 8 intervals
 - If a page history has bits as 11000100 then it has been used more than a page with bits as 01110111
- At regular intervals (e.g., 100 milliseconds) a timer interrupt transfers control to the OS
- OS shifts the reference bit for each page into the high-order bit of its 8-bit byte and shifts the other bits to the right by 1 bit
- Low-order bit falls off and it is discarded
- Page with lowest number is LRU

LRU Page Replacement

- ❑ The last implementation is an approximation of LRU
- ❑ Commonly found

2nd Chance Page Replacement Algorithm

- ❑ A modification to FIFO avoids the problem of throwing out a heavily used page
- ❑ Each page entry has a reference bit, R.
- ❑ The modification is based on an inspection of the R bit of the oldest page
- ❑ If $R=0$ the page is both old and unused
 - Replace immediately
- ❑ If $R=1$
 - R is cleared and put at the end of the list
 - The load time is reset to the current time

Other Algorithms

- ❑ Least frequently used (LFU)
- ❑ Most frequently used (MFU)

- ❑ Most OS's use LRU

Least Recently Used (LRU) Algorithm

- ❑ Why does LRU work?
- ❑ Consider the following code segment:

```
sum = 0;  
for (i=0; i < n; i++)  
{  
    sum = sum + a[i];  
}
```

- ❑ What do we see here?
 - We see that `a[i+1]` is accessed soon after `a[i]`
 - We see that the `sum` is periodically referenced

Least Recently Used (LRU) Algorithm

- ❑ What else do we see?
 - We are cycling through the for-loop repeatedly
- ❑ The program exhibits
 - **Temporal locality**: recently-referenced items are likely to be referenced in the near future
 - Referencing `sum`
 - For loop instructions
 - **Spatial locality**: Items with nearby addresses tend to be referenced close together in time
 - Variable `a[i+1]` accessed after variable `a[i]`

The Working Set Model

- ❑ Processes tend to exhibit a **locality of reference** e.g.,
 - This means that during any phase of execution, the process references only a relatively small fraction of its pages
- ❑ The set of pages that a process is currently using is called its **working set**
- ❑ If the entire working set is in memory, the process will run without causing many faults until it moves into another execution phase
 - Example: Move to another loop

LRU Replacement Algorithm

- ❑ Locality suggests that memory references are on the same set of pages
- ❑ Studies suggest that programs exhibit high spatial and temporal locality

Summary

- ❑ We have studied the need for page replacement algorithms
- ❑ Several algorithms have been discussed including:
 - Optimal
 - FIFO
 - LRU